

# Probabilistic Programming Languages

Guillaume Baudart

*MPRI 2024-2025*

Reminder: Exact Inference

---

Probabilistic Programming Languages

# Example: Coin



Consider a series of coin tosses

- Observations: head or tail
- Each toss is independant
- What is the probability of getting head at the next toss?

Probabilistic model

- Prior:  $z \sim \text{Uniform}(0, 1)$
- Observations: for  $i \in [1, n]$ ,  $x_i \sim \text{Bernoulli}(z)$
- Posterior:  $p(z \mid x_1, \dots, x_n)$ ?

# Example: Coin



Consider a series of coin tosses

- Observations: head or tail
- Each toss is independant
- What is the probability of getting head at the next toss?

Probabilistic model

- Prior:  $z \sim \text{Uniform}(0, 1)$
- Observations: for  $i \in [1, n]$ ,  $x_i \sim \text{Bernoulli}(z)$
- Posterior:  $p(z \mid x_1, \dots, x_n)$ ?

$$\begin{aligned} p(z \mid x_1, \dots, x_n) &= \frac{p(x_1, \dots, x_n \mid z)p(z)}{p(x_1, \dots, x_n)} \\ &= \frac{p(x_1, \dots, x_n \mid z)p(z)}{\int_z p(x_1, \dots, x_n \mid z)} \end{aligned}$$

$$\begin{aligned} p(x_1, \dots, x_n \mid z) &= \prod_{i=1}^n p(x_i \mid z) \\ &= \prod_{i=1}^n z^{x_i} (1 - z)^{1-x_i} \\ &= z^{\sum_{i=1}^n x_i} (1 - z)^{\sum_{i=1}^n (1-x_i)} \\ &= z^{\text{\#heads}} (1 - z)^{\text{\#tails}} \end{aligned}$$

$$\begin{aligned} p(z \mid x_1, \dots, x_n) &= \frac{z^{\text{\#heads}} (1 - z)^{\text{\#tails}}}{\int_z z^{\text{\#heads}} (1 - z)^{\text{\#tails}}} \\ &= \frac{z^{\text{\#heads}} (1 - z)^{\text{\#tails}}}{B(\text{\#heads} + 1, \text{\#tails} + 1)} \\ &= \text{pdf}(\text{Beta}(\text{\#heads} + 1, \text{\#tails} + 1)) \end{aligned}$$

# Example: Coin



Consider a series of coin tosses

- Observations: head or tail
- Each toss is independant
- What is the probability of getting head at the next toss?

Probabilistic model

- Prior:  $z \sim \text{Uniform}(0, 1)$
- Observations: for  $i \in [1, n]$ ,  $x_i \sim \text{Bernoulli}(z)$
- Posterior:  $p(z \mid x_1, \dots, x_n)$ ?

$$\begin{aligned} p(x_1, \dots, x_n \mid z) &= \prod_{i=1}^n p(x_i \mid z) \\ &= \prod_{i=1}^n z^{x_i} (1 - z)^{1-x_i} \\ &= z^{\sum_{i=1}^n x_i} (1 - z)^{\sum_{i=1}^n (1-x_i)} \\ &= z^{\text{\#heads}} (1 - z)^{\text{\#tails}} \end{aligned}$$

$$p(z \mid x_1, \dots, x_n) =$$

=

**What if the model is much more complex?**

**What if we use arbitrary control flow?**

**Can we compute the posterior automatically?**

$$\text{\#heads} (1 - z)^{\text{\#tails}}$$

$$\text{\#heads} (1 - z)^{\text{\#tails}}$$

$$\text{\#heads} (1 - z)^{\text{\#tails}}$$

$$\text{\#heads} + 1, \text{\#tails} + 1)$$

$$\text{beta}(\text{\#heads} + 1, \text{\#tails} + 1))$$

# Semi-Symbolic Inference

---

Probabilistic Programming Languages

# Semi-symbolic inference

Approximate inference can be impractical

- Require lot of computing power
- Can yield poor approximation

Exact inference is often possible

Rao-Blackwellized inference

- Each execution of the model maintains symbolic distributions
- Perform as much exact computation as possible
- Fall back to sampling when symbolic computation fails

Key elements

- Marginalization: remove dependencies
- Conditioning: incorporate observations


# Marginalization

## Remove dependencies

- Parent distribution:  $p(x)$
- Child distribution:  $p(x \mid y)$
- Marginal:  $p(x) = \int p(x \mid y)p(y)$

## Gaussians

- Parent distribution:  $x \sim \mathcal{N}(\mu_0, \sigma_0^2)$
- Child distribution:  $(x \mid y) \sim \mathcal{N}(x, \sigma^2)$
- Marginal:  $y \sim \mathcal{N}(\mu_0, \sigma_0^2 + \sigma^2)$





# Conjugate priors

Bayesian Inference: learn parameters from data

- Latent parameter  $\theta$
- Observed data  $x_1, \dots, x_n$

$$p(\theta \mid x_1, \dots, x_n) = \frac{p(\theta) p(x_1, \dots, x_n \mid \theta)}{p(x_1, \dots, x_n)} \quad (\text{Bayes' theorem})$$

$$\propto p(\theta) p(x_1, \dots, x_n \mid \theta) \quad (\text{Data are constants})$$

If the posterior  $p(\theta \mid x)$  is in the same distribution family as the prior  $p(\theta)$

- Prior and posterior are called *conjugate distributions* w.r.t., the likelihood
- The prior is the *conjugate prior* for the likelihood function

# Conjugate priors examples

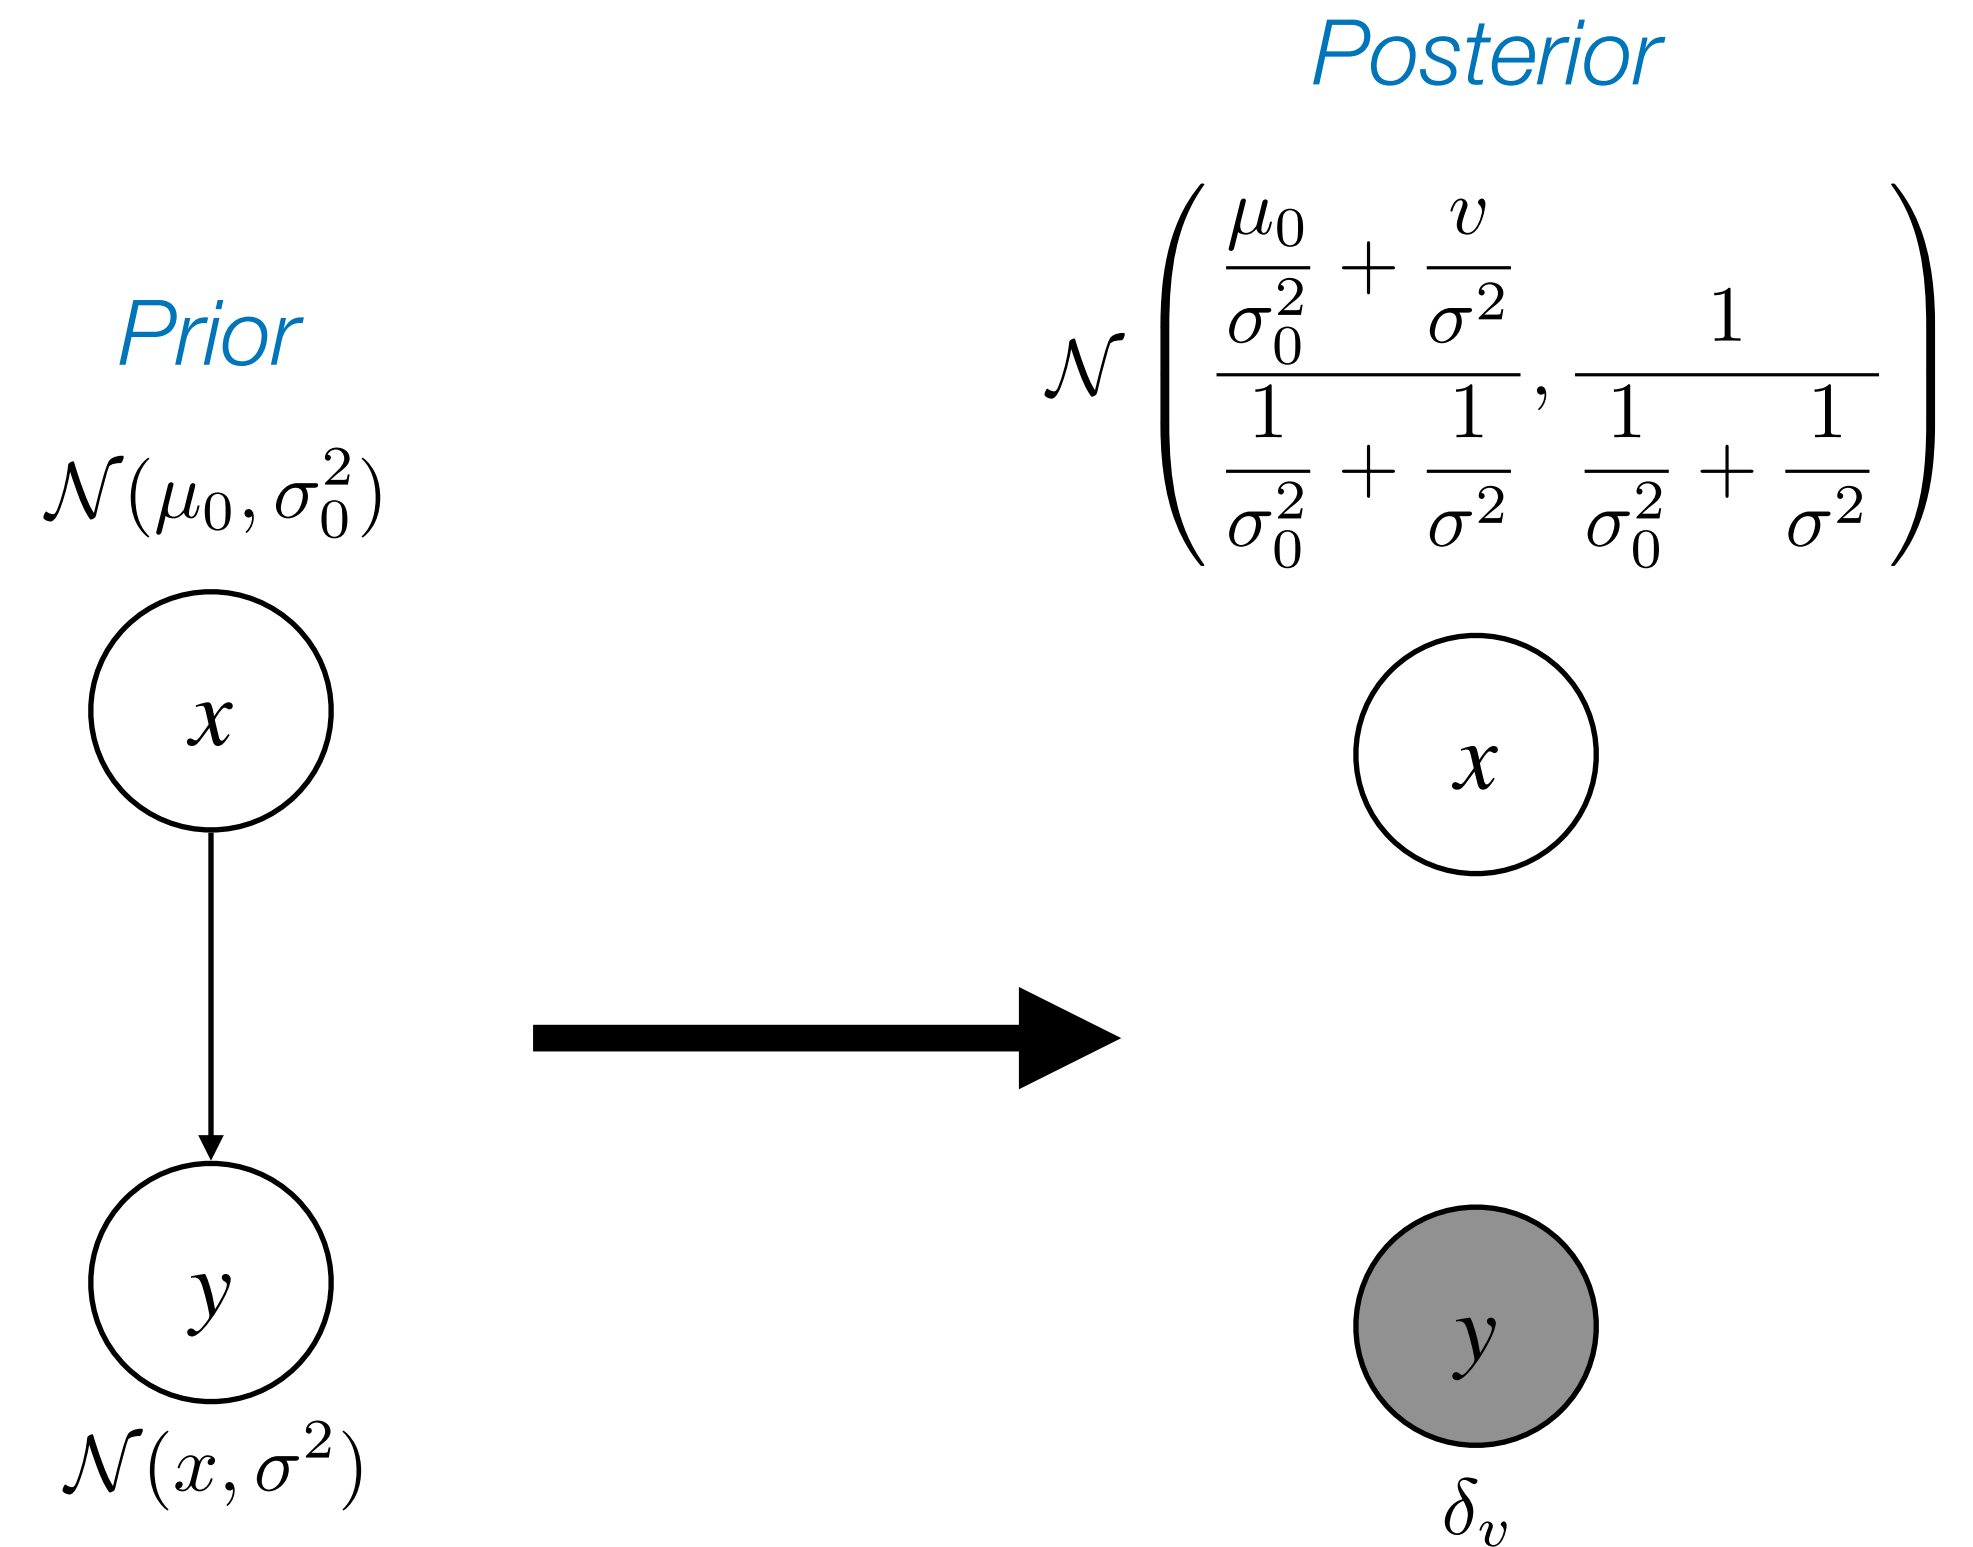
## Beta-Bernoulli

- Prior:  $x \sim \text{Beta}(a, b)$
- Likelihood:  $(y \mid x) \sim \text{Bernoulli}(x)$
- Posterior:  $(x \mid x = v) \sim \text{Beta}(a + v, b + (1 - v))$

## Gaussians

- Prior:  $x \sim \mathcal{N}(\mu_0, \sigma_0^2)$
- Likelihood:  $(y \mid x) \sim \mathcal{N}(x, \sigma^2)$
- Posterior:  $(y \mid x = v) \sim \mathcal{N}(\mu_1, \sigma_1^2)$

$$\begin{aligned}\mu_1 &= \left( \frac{1}{\sigma_0^2} + \frac{1}{\sigma^2} \right)^{-1} \left( \frac{\mu_0}{\sigma_0^2} + \frac{v}{\sigma^2} \right) \\ \sigma_1^2 &= \left( \frac{1}{\sigma_0^2} + \frac{1}{\sigma^2} \right)^{-1}\end{aligned}$$



Try it in BYO-PPL!

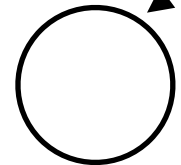
# Delayed Sampling


---

Probabilistic Programming Languages

# Delayed sampling


```
let gauss () =  
  let a = sample gaussian 0. 1. in  
  observe (gaussian z 1.) 2.;  
  observe (gaussian z 1.) 2.;  
  observe (gaussian z 1.) 2.;  
  a
```

$\mathcal{N}(0, 1)$  




# Delayed sampling

```
let gauss () =  
  let a = sample gaussian 0. 1. in  
  observe (gaussian z 1.) 2.;  
  observe (gaussian z 1.) 2.;  
  observe (gaussian z 1.) 2.;  
  a
```



# Delayed sampling

```
let gauss () =  
  let a = sample gaussian 0. 1. in  
  observe (gaussian z 1.) 2.;  
  observe (gaussian z 1.) 2.;  
  observe (gaussian z 1.) 2.;  
  a
```




## Conditionning

### Gaussians

- Prior:  $x \sim \mathcal{N}(\mu_0, \sigma_0^2)$
- Likelihood:  $(y \mid x) \sim \mathcal{N}(x, \sigma^2)$
- Posterior:  $(y \mid x = v) \sim \mathcal{N}(\mu_1, \sigma_1^2)$

$$\mu_1 = \left( \frac{1}{\sigma_0^2} + \frac{1}{\sigma^2} \right)^{-1} \left( \frac{\mu_0}{\sigma_0^2} + \frac{v}{\sigma^2} \right)$$
$$\sigma_1^2 = \left( \frac{1}{\sigma_0^2} + \frac{1}{\sigma^2} \right)^{-1}$$

# Delayed sampling



```
let gauss () =  
  let a = sample gaussian 0. 1. in  
  observe (gaussian z 1.) 2.;  
  observe (gaussian z 1.) 2.;  
  a
```

## Conditionning


### Gaussians

- Prior:  $x \sim \mathcal{N}(\mu_0, \sigma_0^2)$
- Likelihood:  $(y | x) \sim \mathcal{N}(x, \sigma^2)$
- Posterior:  $(y | x = v) \sim \mathcal{N}(\mu_1, \sigma_1^2)$

$$\mu_1 = \left(\frac{1}{\sigma_0^2} + \frac{1}{\sigma^2}\right)^{-1} \left(\frac{\mu_0}{\sigma_0^2} + \frac{v}{\sigma^2}\right)$$
$$\sigma_1^2 = \left(\frac{1}{\sigma_0^2} + \frac{1}{\sigma^2}\right)^{-1}$$



# Delayed sampling



```

let gauss () =
  let a = sample gaussian 0. 1. in
  observe (gaussian z 1.) 2.;
  observe (gaussian z 1.) 2.;
  observe (gaussian z 1.) 2.;
  a
  
```

## Conditionning

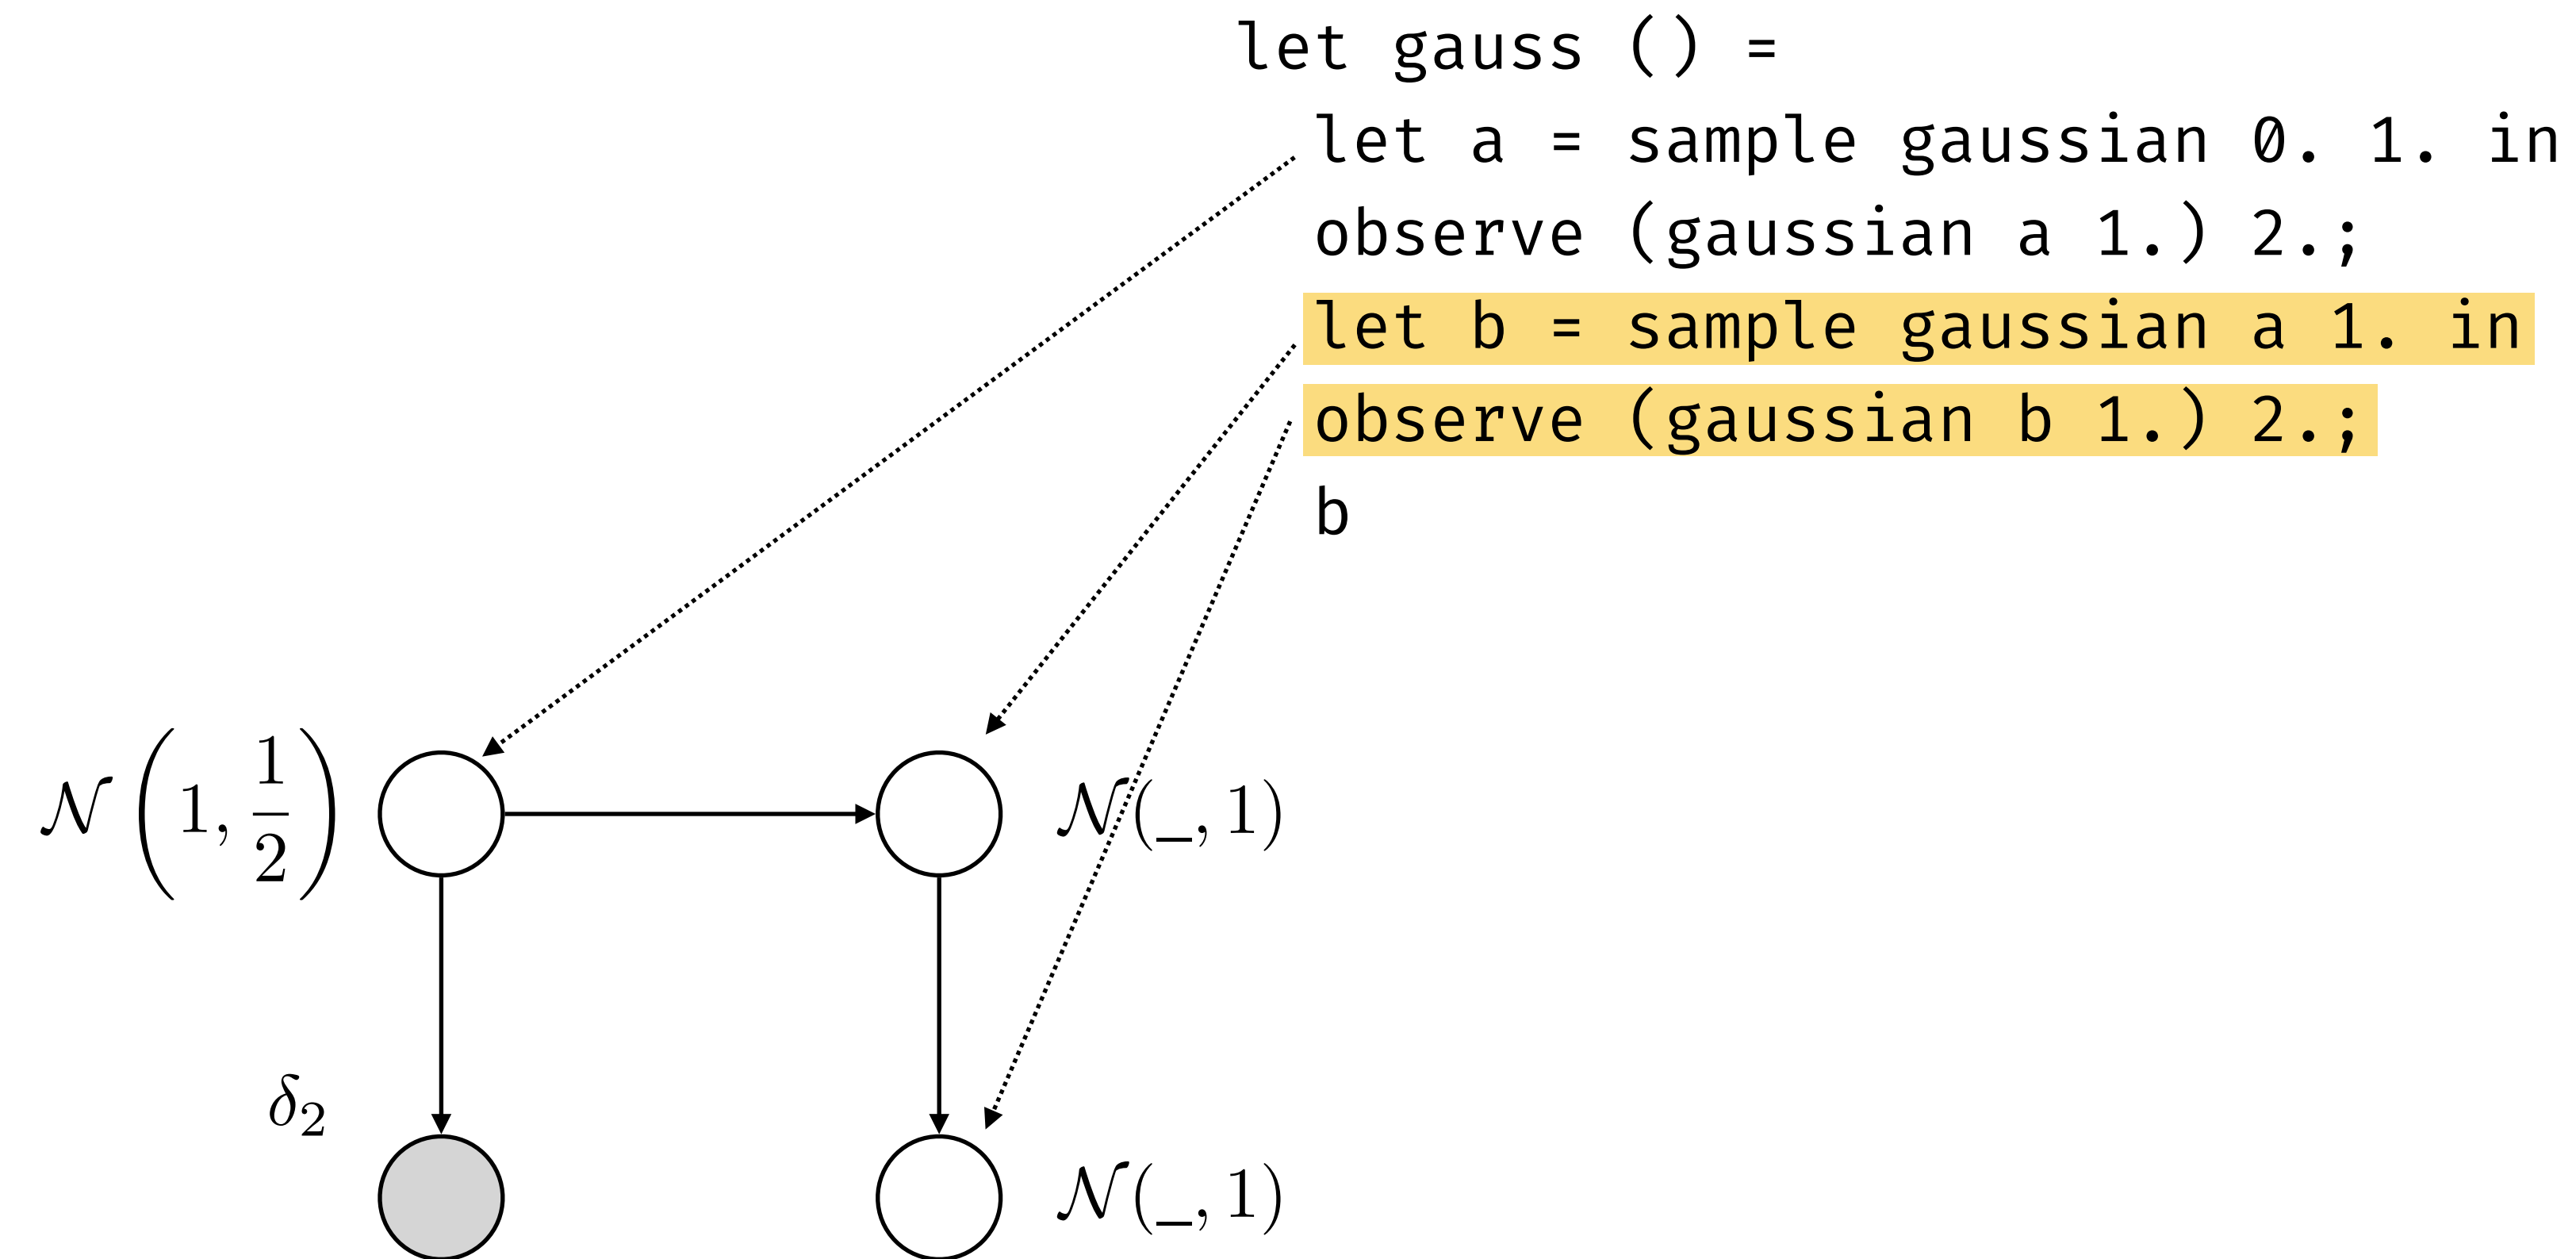
### Gaussians

- Prior:  $x \sim \mathcal{N}(\mu_0, \sigma_0^2)$
- Likelihood:  $(y | x) \sim \mathcal{N}(x, \sigma^2)$
- Posterior:  $(y | x = v) \sim \mathcal{N}(\mu_1, \sigma_1^2)$

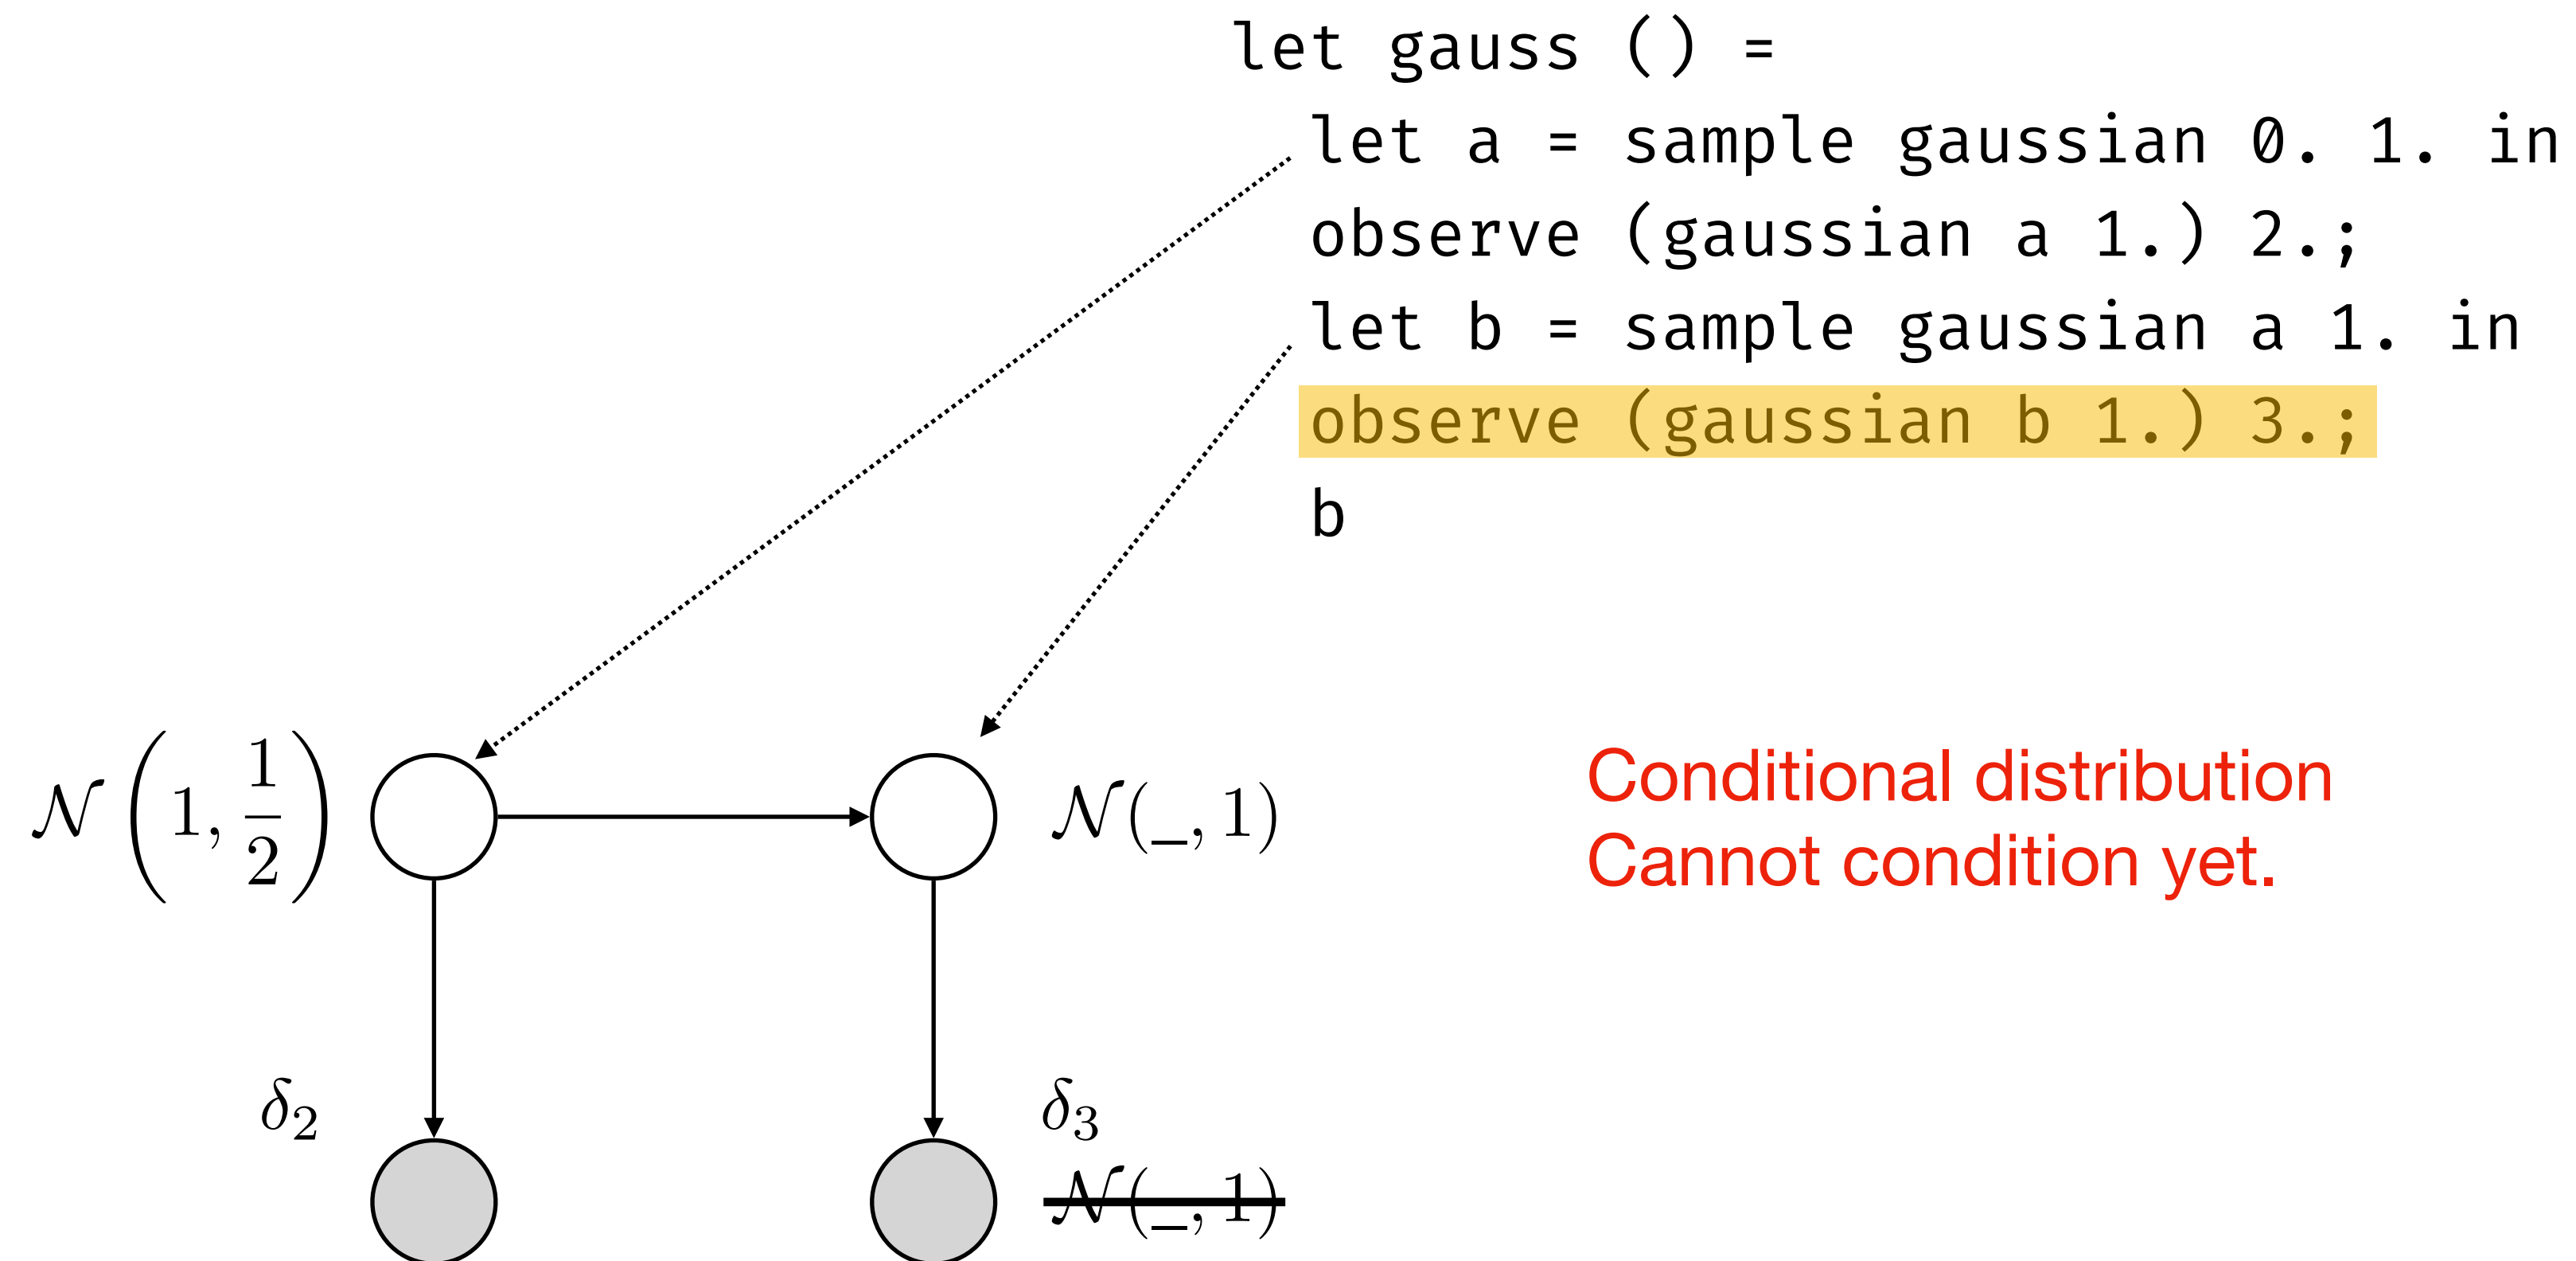
$$\mu_1 = \left(\frac{1}{\sigma_0^2} + \frac{1}{\sigma^2}\right)^{-1} \left(\frac{\mu_0}{\sigma_0^2} + \frac{v}{\sigma^2}\right)$$

$$\sigma_1^2 = \left(\frac{1}{\sigma_0^2} + \frac{1}{\sigma^2}\right)^{-1}$$

# Delayed sampling




# Delayed sampling



# Delayed sampling

```
let gauss () =  
  let a = sample gaussian 0. 1. in  
  observe (gaussian a 1.) 2.;  
  let b = sample gaussian a 1. in  
  observe (gaussian b 1.) 3.
```




## Marginalization

### Gaussians

- Parent distribution:  $x \sim \mathcal{N}(\mu_0, \sigma_0^2)$
- Child distribution:  $(x | y) \sim \mathcal{N}(x, \sigma^2)$
- Marginal:  $y \sim \mathcal{N}(\mu_0, \sigma_0^2 + \sigma^2)$

# Delayed sampling

```
let gauss () =
  let a = sample gaussian 0. 1. in
  observe (gaussian a 1.) 2.;
  let b = sample gaussian a 1. in
  observe (gaussian b 1.) 3.
```



## Conditioning

### Gaussians

- Prior:  $x \sim \mathcal{N}(\mu_0, \sigma_0^2)$
- Likelihood:  $(y | x) \sim \mathcal{N}(x, \sigma^2)$
- Posterior:  $(y | x = v) \sim \mathcal{N}(\mu_1, \sigma_1^2)$

$$\mu_1 = \left(\frac{1}{\sigma_0^2} + \frac{1}{\sigma^2}\right)^{-1} \left(\frac{\mu_0}{\sigma_0^2} + \frac{v}{\sigma^2}\right)$$

$$\sigma_1^2 = \left(\frac{1}{\sigma_0^2} + \frac{1}{\sigma^2}\right)^{-1}$$

# Delayed sampling: summary

## `infer` model data

- Launch  $n$  independent executions (approximate inference)
- Maintain a graph (Bayesian model for the current execution)
- Returns a mixture of symbolic distributions

## `sample` d

- add a node to the graph (random variable)
- keep track of dependencies

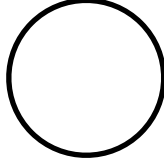
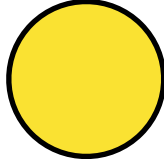
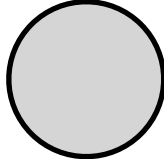
## `observe`

- add a node to the graph (observed random variable)
- marginalize parent (recursively until the root)
- condition the parent on the observation


## If error (conditioning or marginalization fails)

- Sample parent
- Update the children


## Node state a runtime

-  Initialized (conditional distribution)
-  Marginalized (marginal distribution)
-  Realized (observed value)

# Delayed sampling: actual sampling



no conjugacy



single m-path

Delayed sampling is limited to forests.  
Can we do better?

# Semi-symbolic inference

---

## Probabilistic Programming Languages



# Semi-symbolic inference

## Gaussians



- Prior:  $x \sim \mathcal{N}(\mu_0, \sigma_0^2)$
- Likelihood:  $(y \mid x) \sim \mathcal{N}(x, \sigma^2)$
- Posterior:  $(y \mid x = v) \sim \mathcal{N}(\mu_1, \sigma_1^2)$

$$\mu_1 = \left( \frac{1}{\sigma_0^2} + \frac{1}{\sigma^2} \right)^{-1} \left( \frac{\mu_0}{\sigma_0^2} + \frac{v}{\sigma^2} \right)$$

$$\sigma_1^2 = \left( \frac{1}{\sigma_0^2} + \frac{1}{\sigma^2} \right)^{-1}$$

*Prior*

$\mathcal{N}(\mu_0, \sigma_0^2)$





$\mathcal{N}(x, \sigma^2)$

swap

*Posterior*

$$\mathcal{N} \left( \frac{\frac{\mu_0}{\sigma_0^2} + \frac{y}{\sigma^2}}{\frac{1}{\sigma_0^2} + \frac{1}{\sigma^2}}, \frac{1}{\frac{1}{\sigma_0^2} + \frac{1}{\sigma^2}} \right)$$



$\mathcal{N}(\mu_0, \sigma_0^2 + \sigma^2)$

*Marginal*


# Semi-symbolic inference

## Main idea

- Maintain symbolic expressions involving random variables
- Each operation triggers a sequence of swaps
- Swap change the dependency between two variables but preserves the symbolic state
- Only root nodes (with no parent) can be observed



## Closed families

- A swap always succeed and the resulting state remains in the family.
- E.g., linear Gaussians models, and finite discrete models.



# Semi-symbolic inference

```
let wheels(l_rate, r_rate) =  
  let omega = sample (gaussian 0. 50.0) in  
  let vel = sample (gaussian 0. 50.0) in  
  observe (gaussian(vel - 2.0 * omega, 1.0) l_rate;  
  observe (gaussian(vel + 2.0 * omega, 1.0) r_rate;  
  omega, vel
```





# Semi-symbolic inference

```

let wheels(l_rate, r_rate) =
  let omega = sample (gaussian 0. 50.0) in
  let vel = sample (gaussian 0. 50.0) in
  observe (gaussian(vel - 2.0 * omega, 1.0) l_rate;
  observe (gaussian(vel + 2.0 * omega, 1.0) r_rate;
  omega, vel

```





swap( $l$ ,  $v$ )

# Semi-symbolic inference

```

let wheels(l_rate, r_rate) =
  let omega = sample (gaussian 0. 50.0) in
  let vel = sample (gaussian 0. 50.0) in
  observe (gaussian(vel - 2.0 * omega, 1.0) l_rate;
  observe (gaussian(vel + 2.0 * omega, 1.0) r_rate;
  omega, vel

```



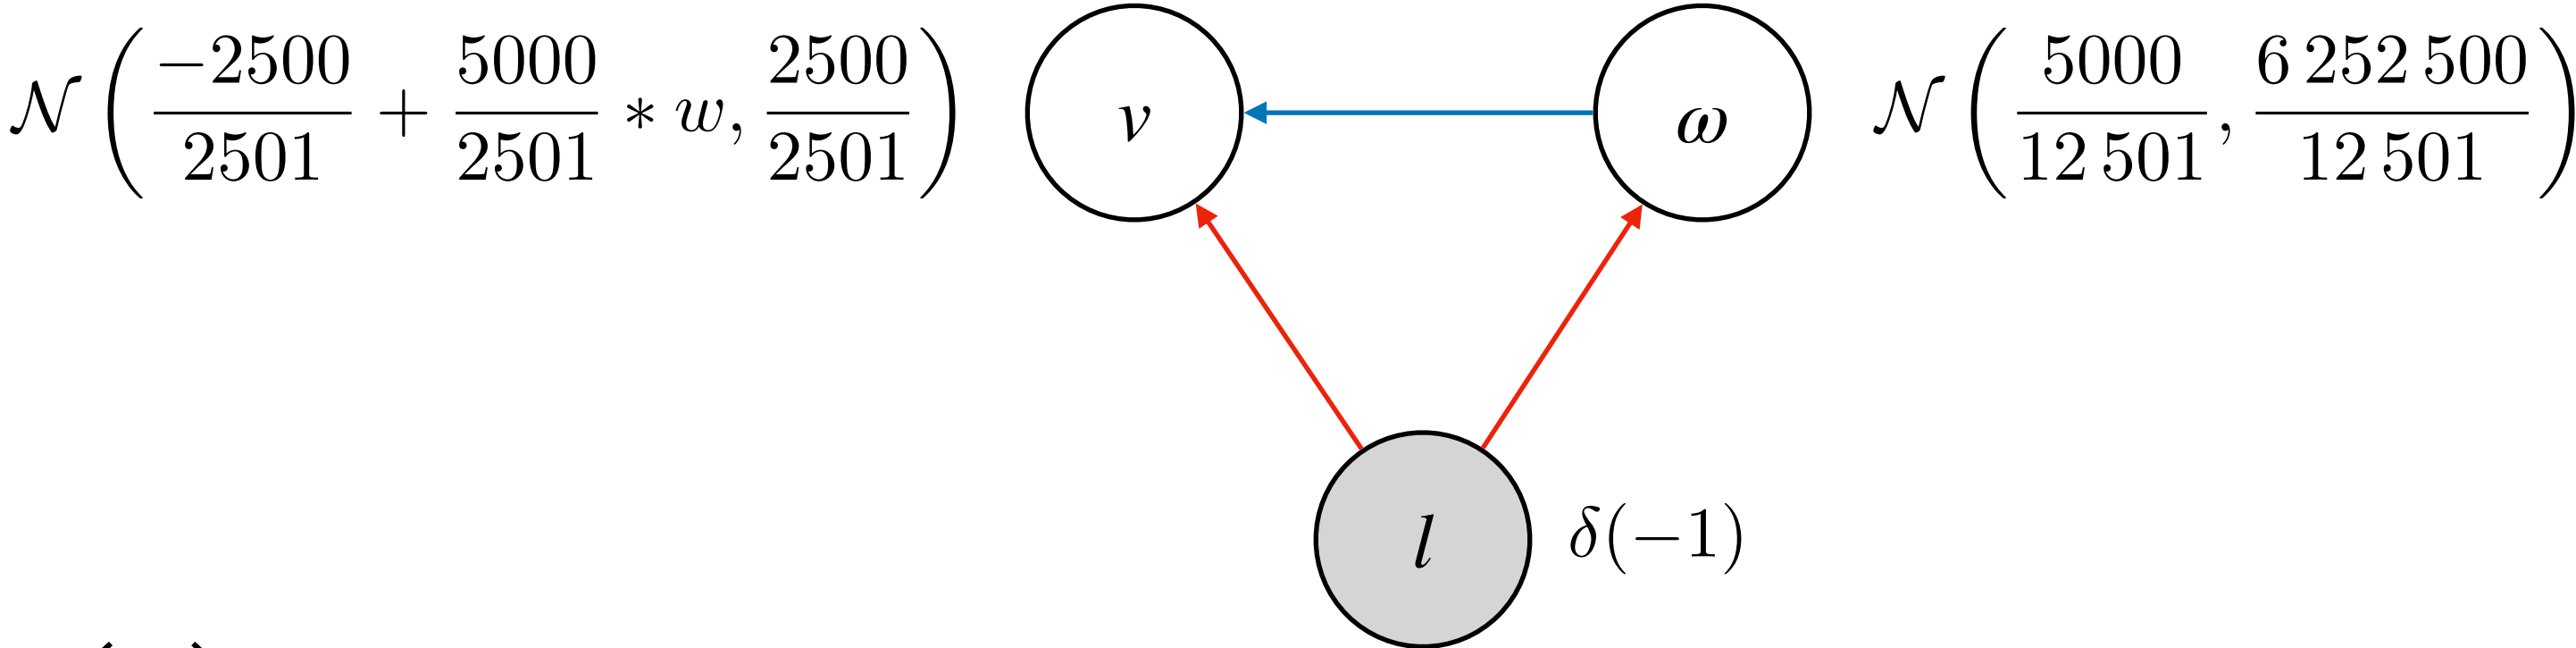

swap(l, w)

# Semi-symbolic inference

```

let wheels(l_rate, r_rate) =
  let omega = sample (gaussian 0. 50.0) in
  let vel = sample (gaussian 0. 50.0) in
  observe (gaussian(vel - 2.0 * omega, 1.0) l_rate;
  observe (gaussian(vel + 2.0 * omega, 1.0) r_rate;
  omega, vel

```



value(l)

# Semi-symbolic: summary

## `infer` model data

- Launch  $n$  independent executions (approximate inference)
- Maintain a graph (Bayesian model for the current execution)
- Returns a mixture of symbolic distributions

## `sample` d

- add a node to the graph (random variable)
- keep track of dependencies

## `observe`

- add a node to the graph (observed random variable)
- swaps to turn the node into a root
- condition and simplify the graph

## If error (swap fails)

- Sample (some) parents
- Update the children

More general than delayed sampling

# Reactive Probabilistic Programming

---

Probabilistic Programming Languages



# Uncertainty in embedded systems

## Synchronous languages

- High-level specification language
- Generate correct-by-construction embedded code
- Industrial tool: ANSYS Scade

## Challenges


- Noisy environment, perceived through noisy sensors
- Interaction with other autonomous entities

## Existing approaches

- Manually implement stochastic controller: *Can be error prone*
- Offline statistical tests: *Requires up-to-date offline data*

## Reactive Probabilistic Programming

- Synchronous languages with probabilistic constructs
- Make the probabilistic model explicit
- Automatically learn posterior distributions from observations

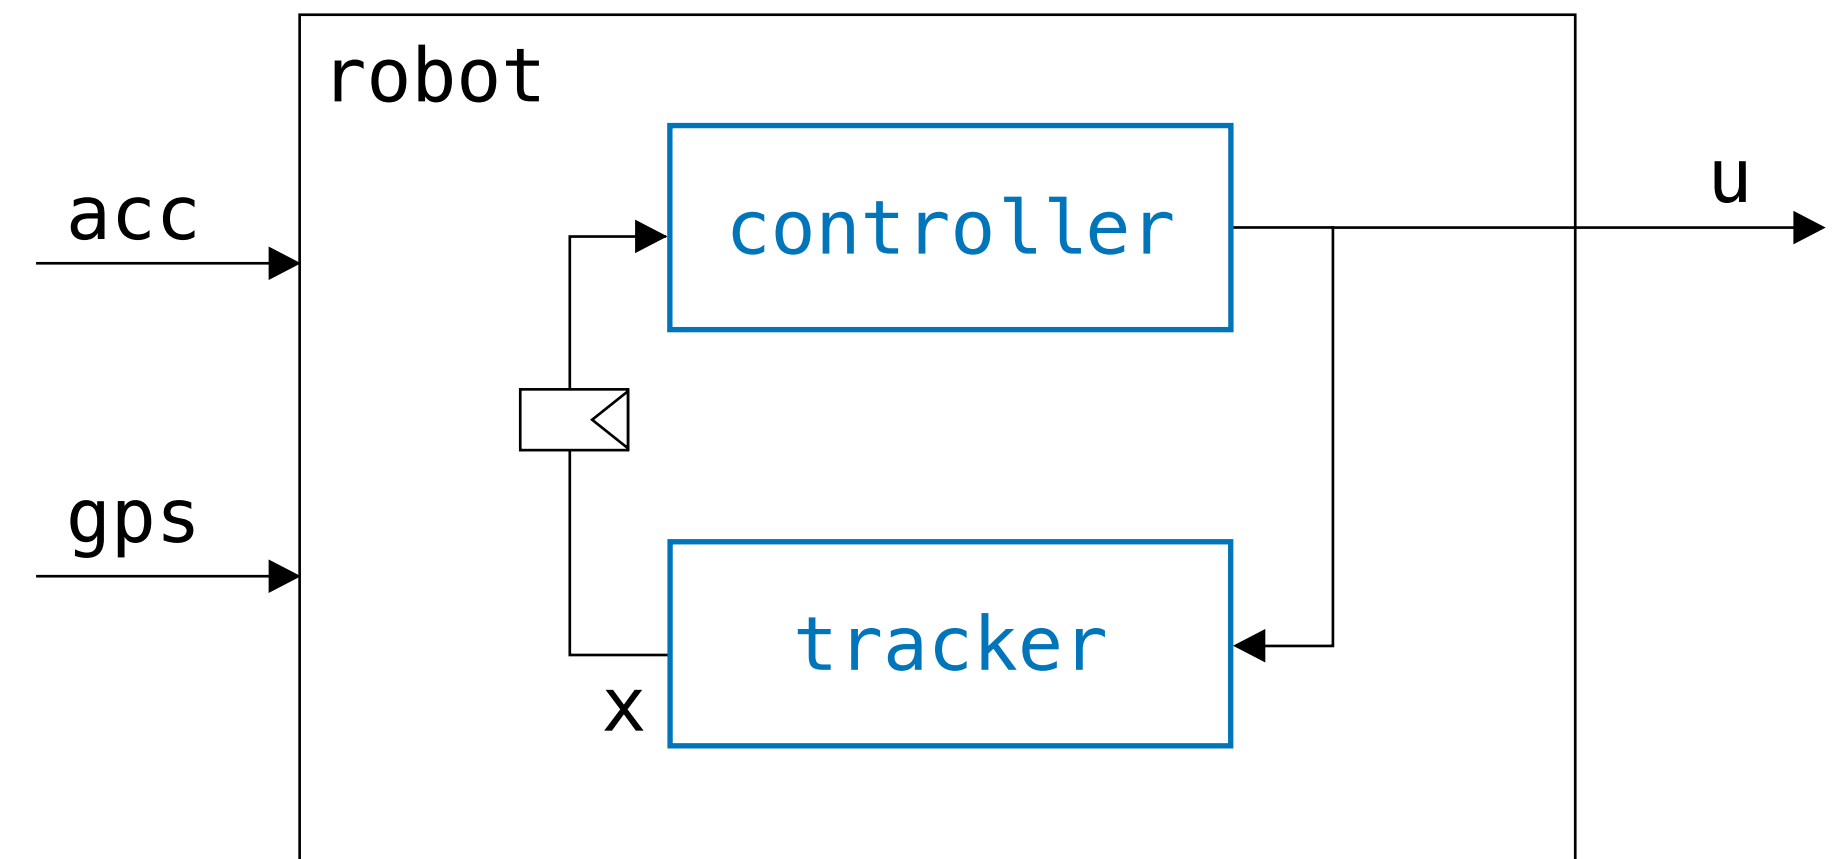


# Reactive systems

Synchronous data-flow languages and block diagrams

- Signal: stream of values
- System: stream processor

State:  $x : (\textit{position} \times \textit{velocity} \times \textit{acceleration})$



# Reactive probabilistic systems

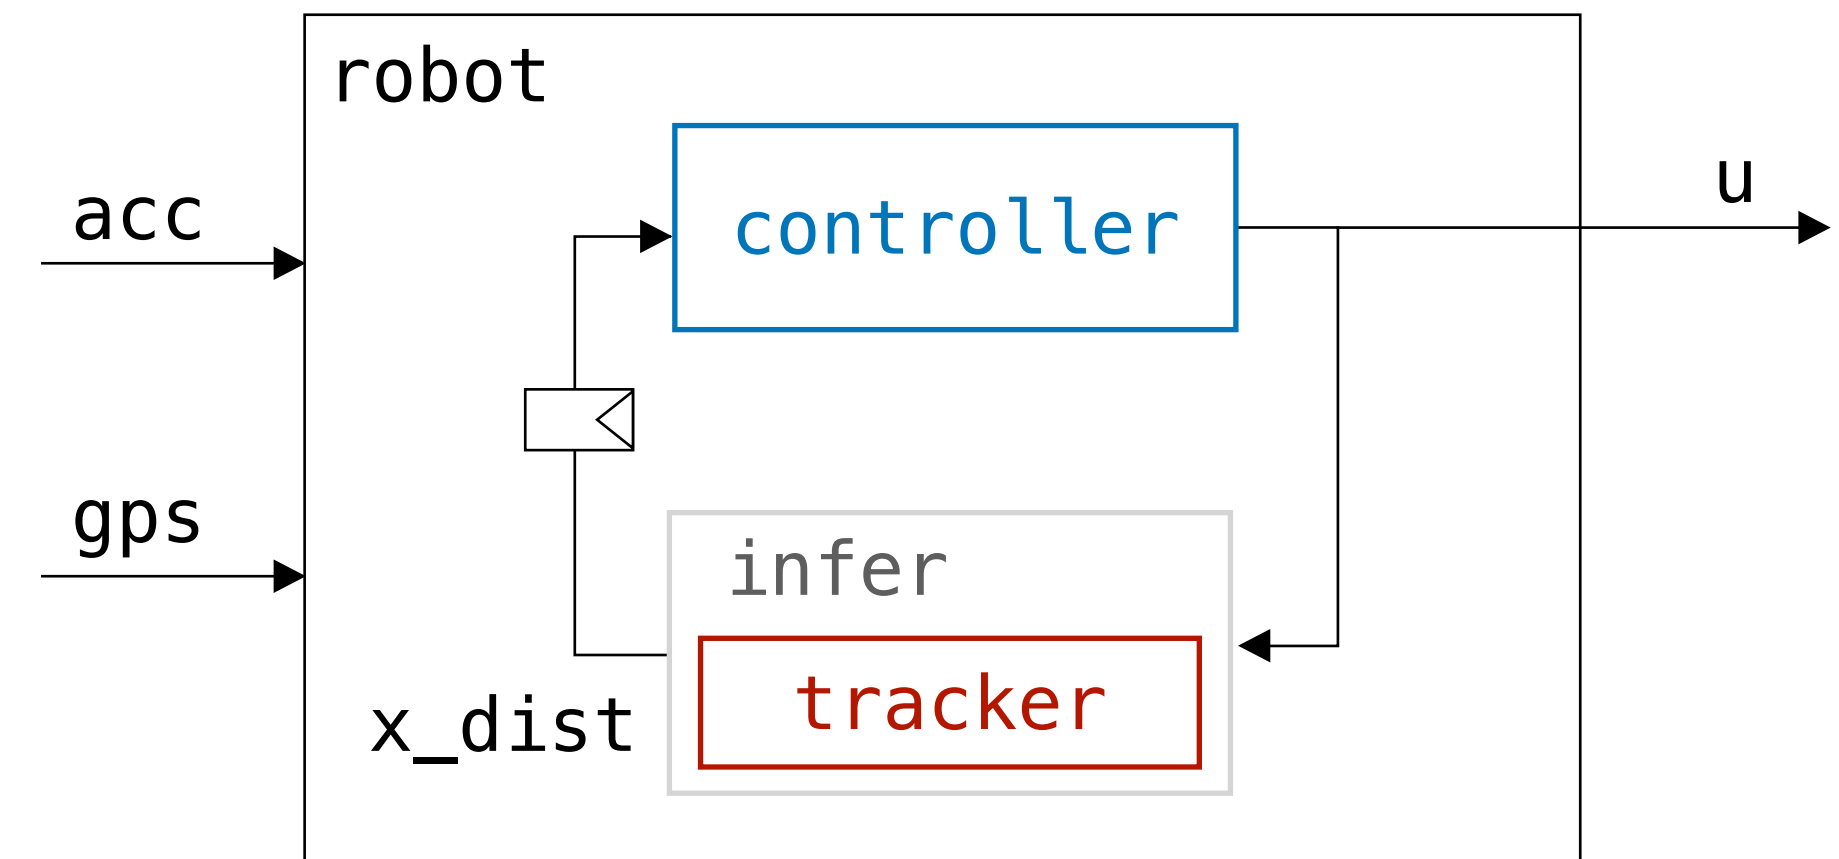
## Synchronous data-flow languages and block diagrams

- Signal: stream of values
- System: stream processor

## ProbZelus: add support to deal with uncertainty

- Extend a synchronous language
- Parallel composition: deterministic/probabilistic
- Inference-in-the-loop
- Streaming inference

State:  $x\_dist: (position \times velocity \times acceleration) dist$



# Synchronous Programming

---

Reactive Probabilistic Programming

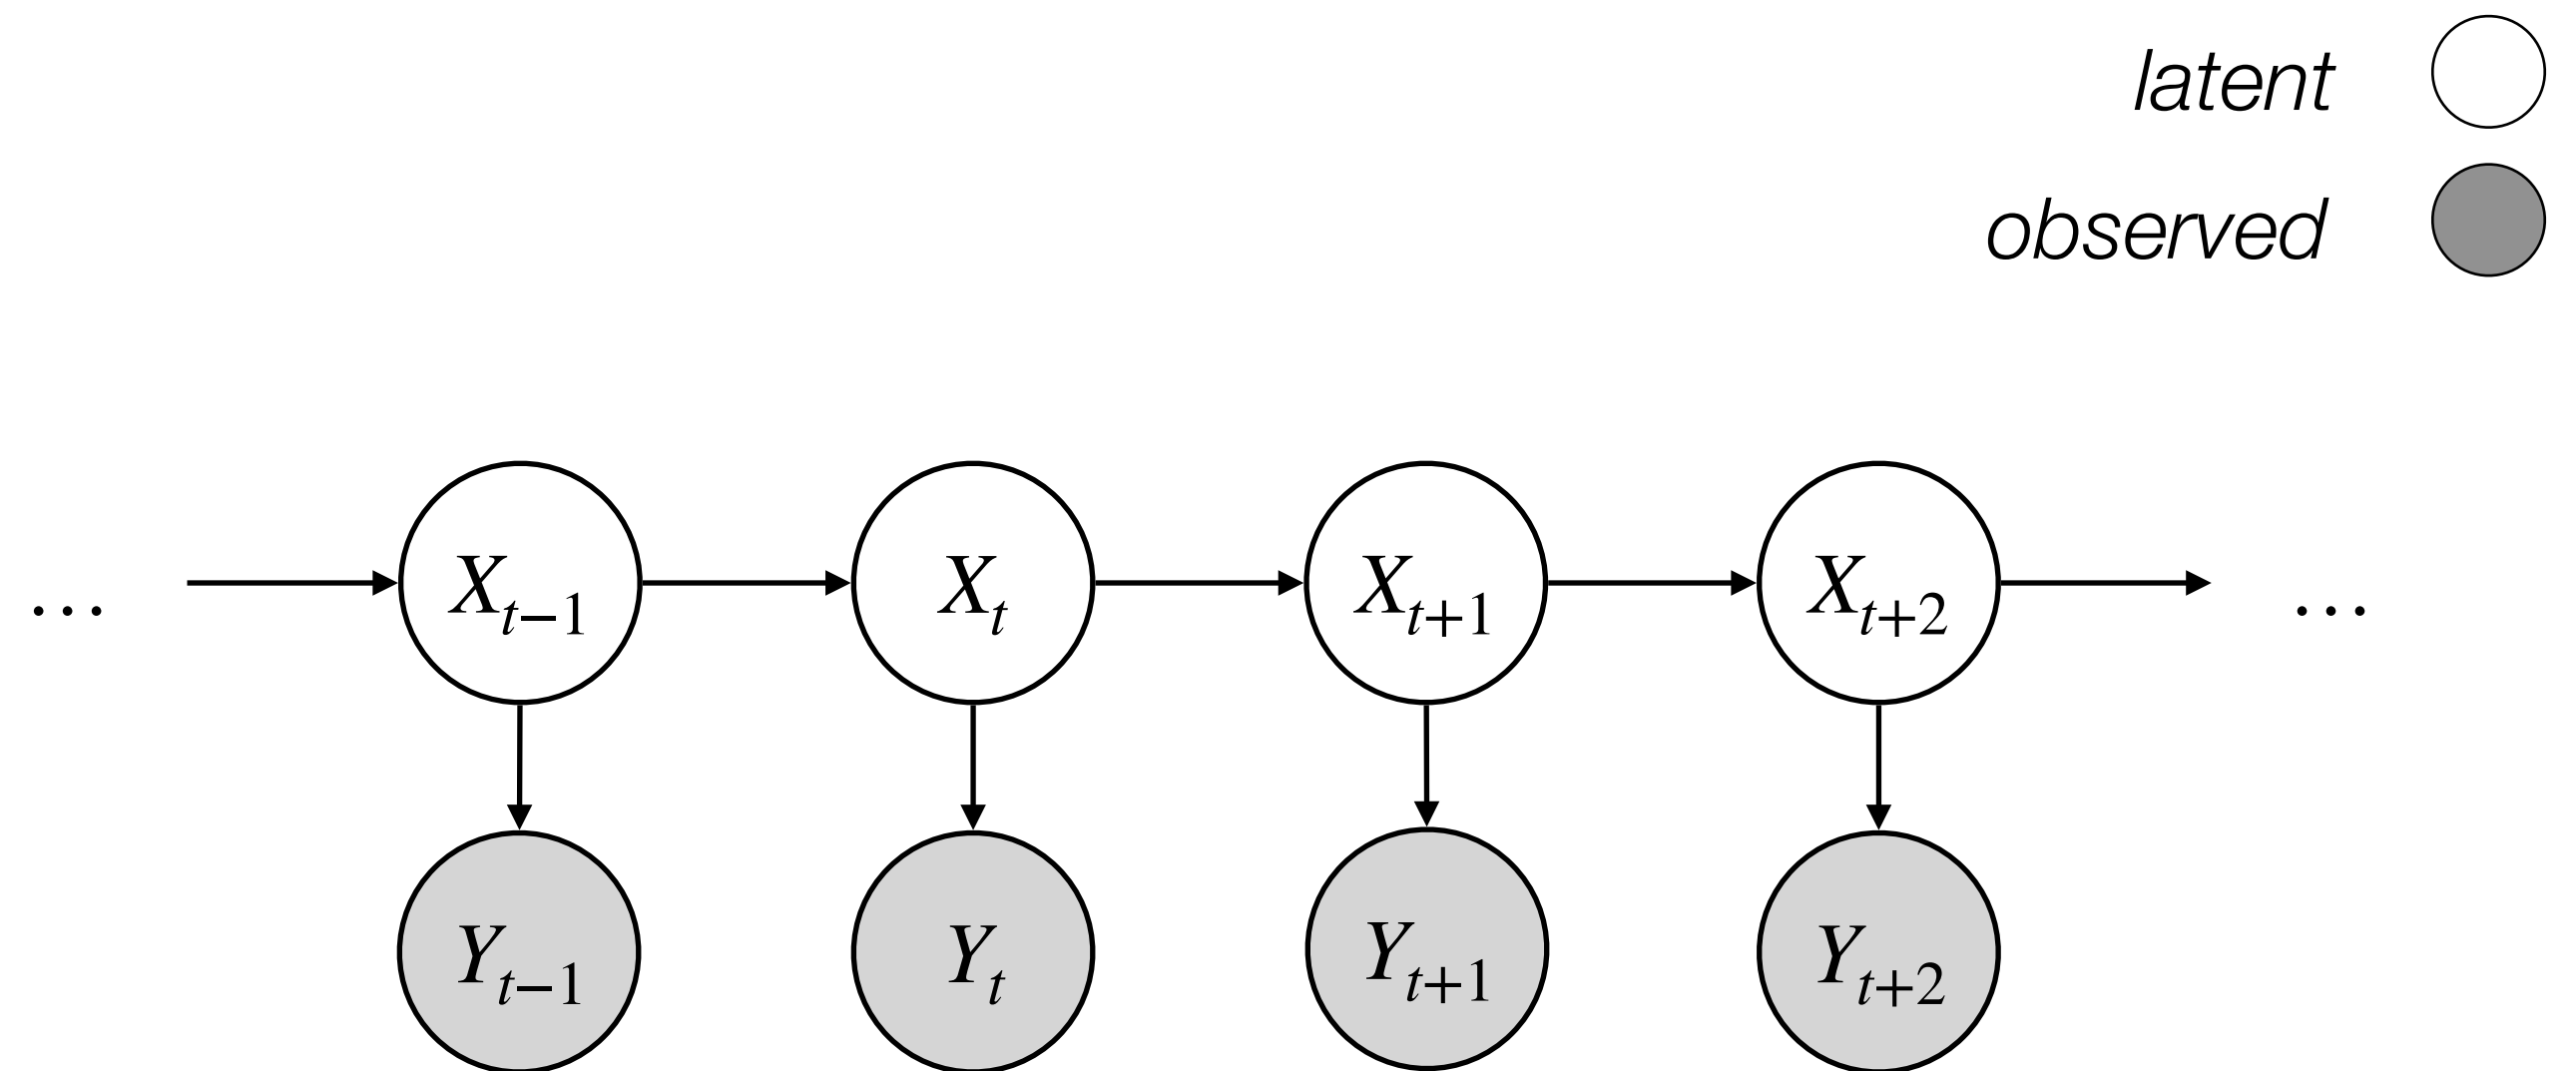
# Example: tracker

## Model

- Linear motion:  $X_k \sim \mathcal{N}(FX_{k-1}, Q)$
- Observation:  $Y_k \sim \mathcal{N}(HX_k, R)$

E.g., with  $Q$  and  $R$  constant noise matrices

- $X_k = \begin{pmatrix} p_k \\ v_k \end{pmatrix}$  (position, velocity)
- $F = \begin{pmatrix} 1 & dt \\ 0 & 1 \end{pmatrix}$  (discrete integration)
- $H = \begin{pmatrix} 1 \\ 0 \end{pmatrix}$  (projection)



$$\begin{aligned} X_k^{\text{pred}} &= FX_{k-1}^{\text{est}} \\ X_k^{\text{est}} &= X_k^{\text{pred}} + K_k(Y_k - HX_k^{\text{pred}}) \\ S_k &= (R + HP_k^{\text{pred}}H^T)^{-1} \\ K_k &= P_k^{\text{pred}}H^TS_k \\ P_k^{\text{pred}} &= Q + FP_{k-1}^{\text{est}}F^T \\ P_k^{\text{est}} &= P_k^{\text{pred}} - K_kHP_k^{\text{pred}} \end{aligned}$$

**Solution: Kalman filter**

# Reactive synchronous programming

## Dataflow synchronous programming

- Set of stream equations
- Discrete logical time steps
- At each step, compute the current value given inputs and previous values

```
let node kalman(y) = x_est where
  rec x_pred = f * (x0 → pre x_est)
  and x_est   = x_pred + k * (y - h * x_pred)
  and s       = r + h * p_pred * (transpose h)
  and k       = p_pred * (transpose h) * (inv s)
  and p_pred  = q + f * (p0 → pre p_est) * (transpose f)
  and p_est   = p_pred - k * h * p_pred
```

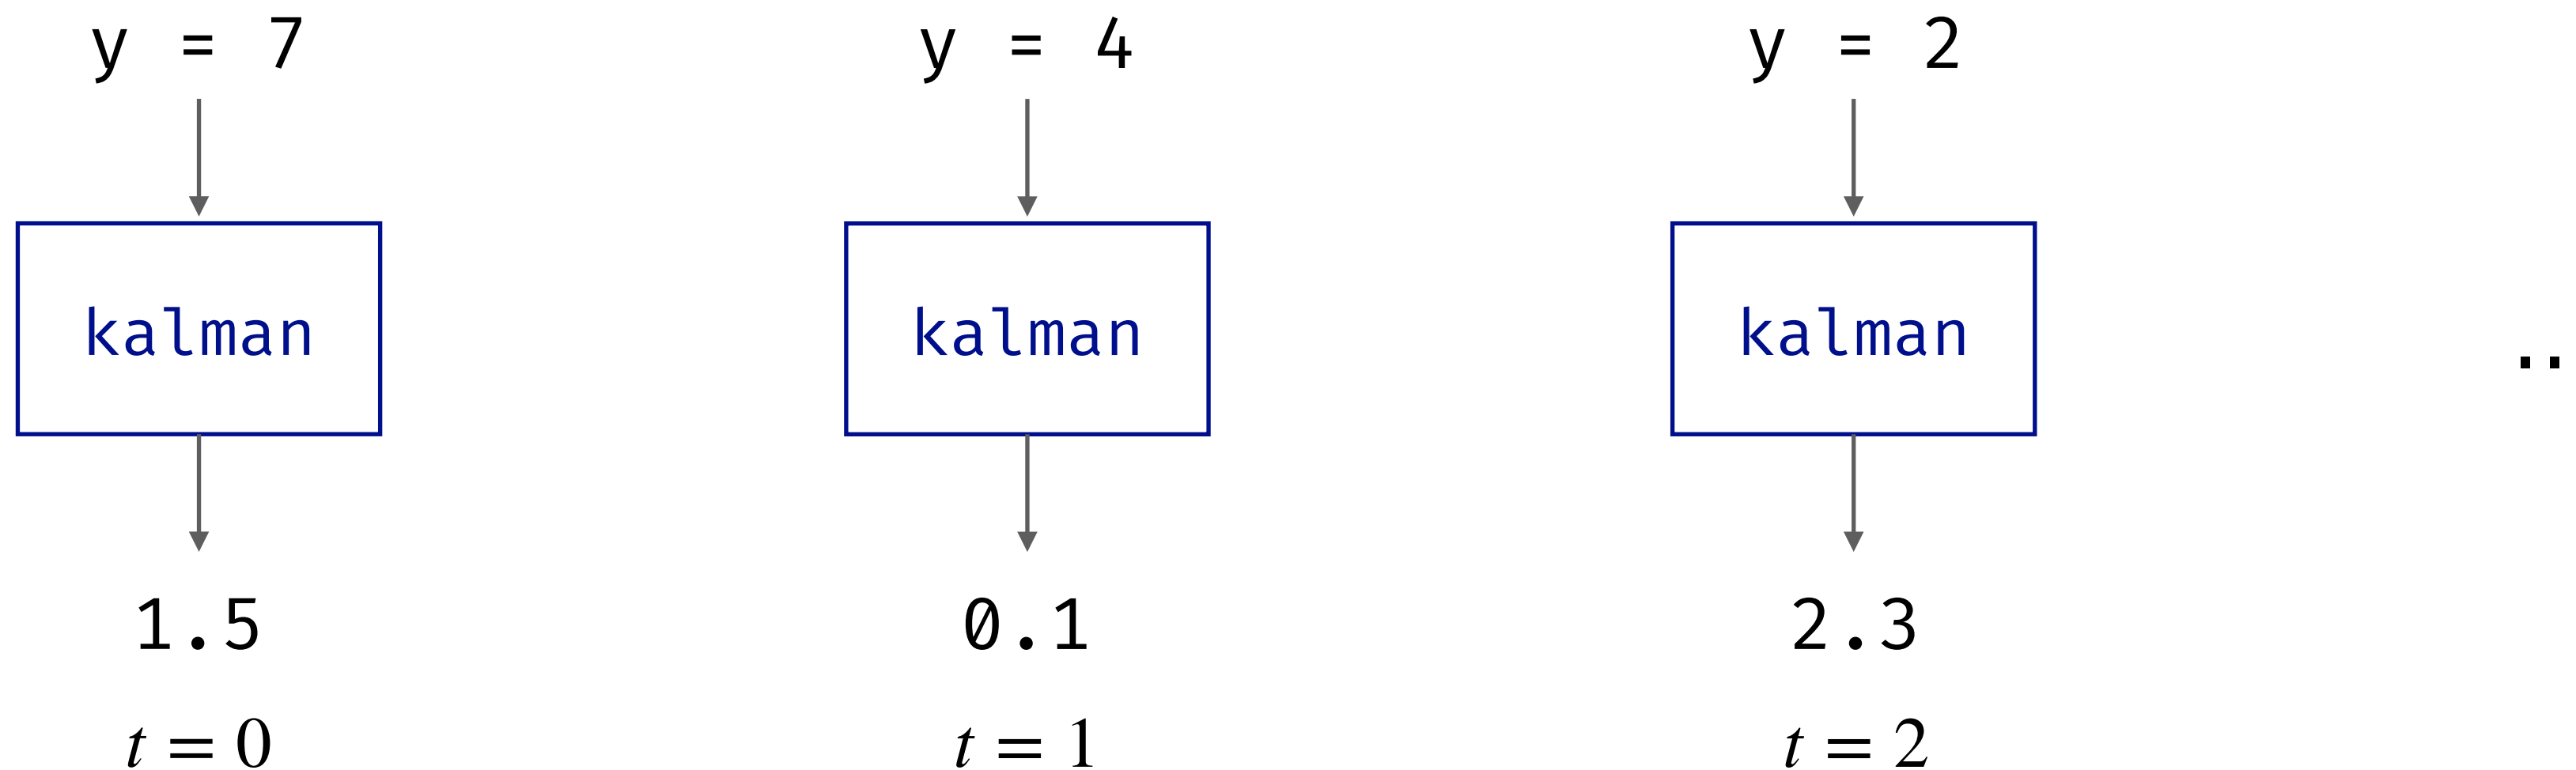
$$\begin{aligned}X_k^{\text{pred}} &= F X_{k-1}^{\text{est}} \\X_k^{\text{est}} &= X_k^{\text{pred}} + K_k (Y_k - H X_k^{\text{pred}}) \\S_k &= R + H P_k^{\text{pred}} H^T \\K_k &= P_k^{\text{pred}} H^T S_k^{-1} \\P_k^{\text{pred}} &= Q + F P_{k-1}^{\text{est}} F^T \\P_k^{\text{est}} &= P_k^{\text{pred}} - K_k H P_k^{\text{pred}}\end{aligned}$$

Solution: Kalman filter

# Reactive synchronous programming

```
let node kalman(y) = x_est where
  rec x_pred = f * (x0 → pre x_est)
  and x_est   = x_pred + k * (y - h * x_pred)
  and s       = r + h * p_pred * (transpose h)
  and k       = p_pred * (transpose h) * (inv s)
  and p_pred  = q + f * (p0 → pre p_est) * (transpose f)
  and p_est   = p_pred - k * h * p_pred
```

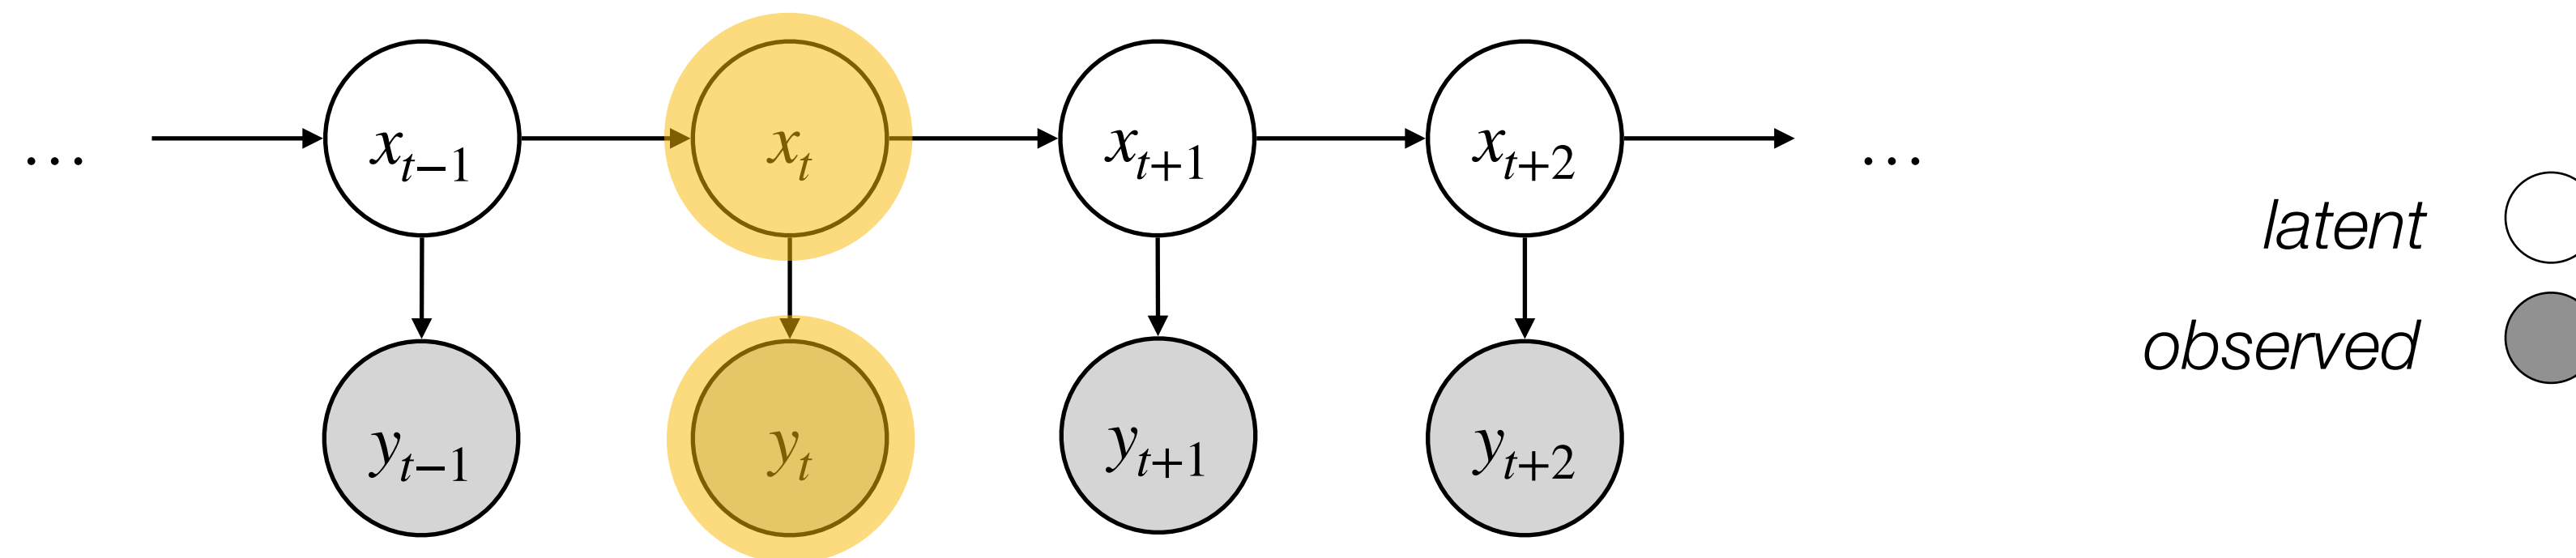
What if the assumptions change?  
What if the model is not linear?



# Reactive probabilistic programming

## Probabilistic constructs

- $x = \text{sample}(d)$ : introduce a random variable  $x$  of distribution  $d$
- $\text{observe}(d, y)$ : condition on the fact that  $y$  was sampled from  $d$
- $\text{infer } m \ y$ : compute posterior distribution of  $m$  given  $y$



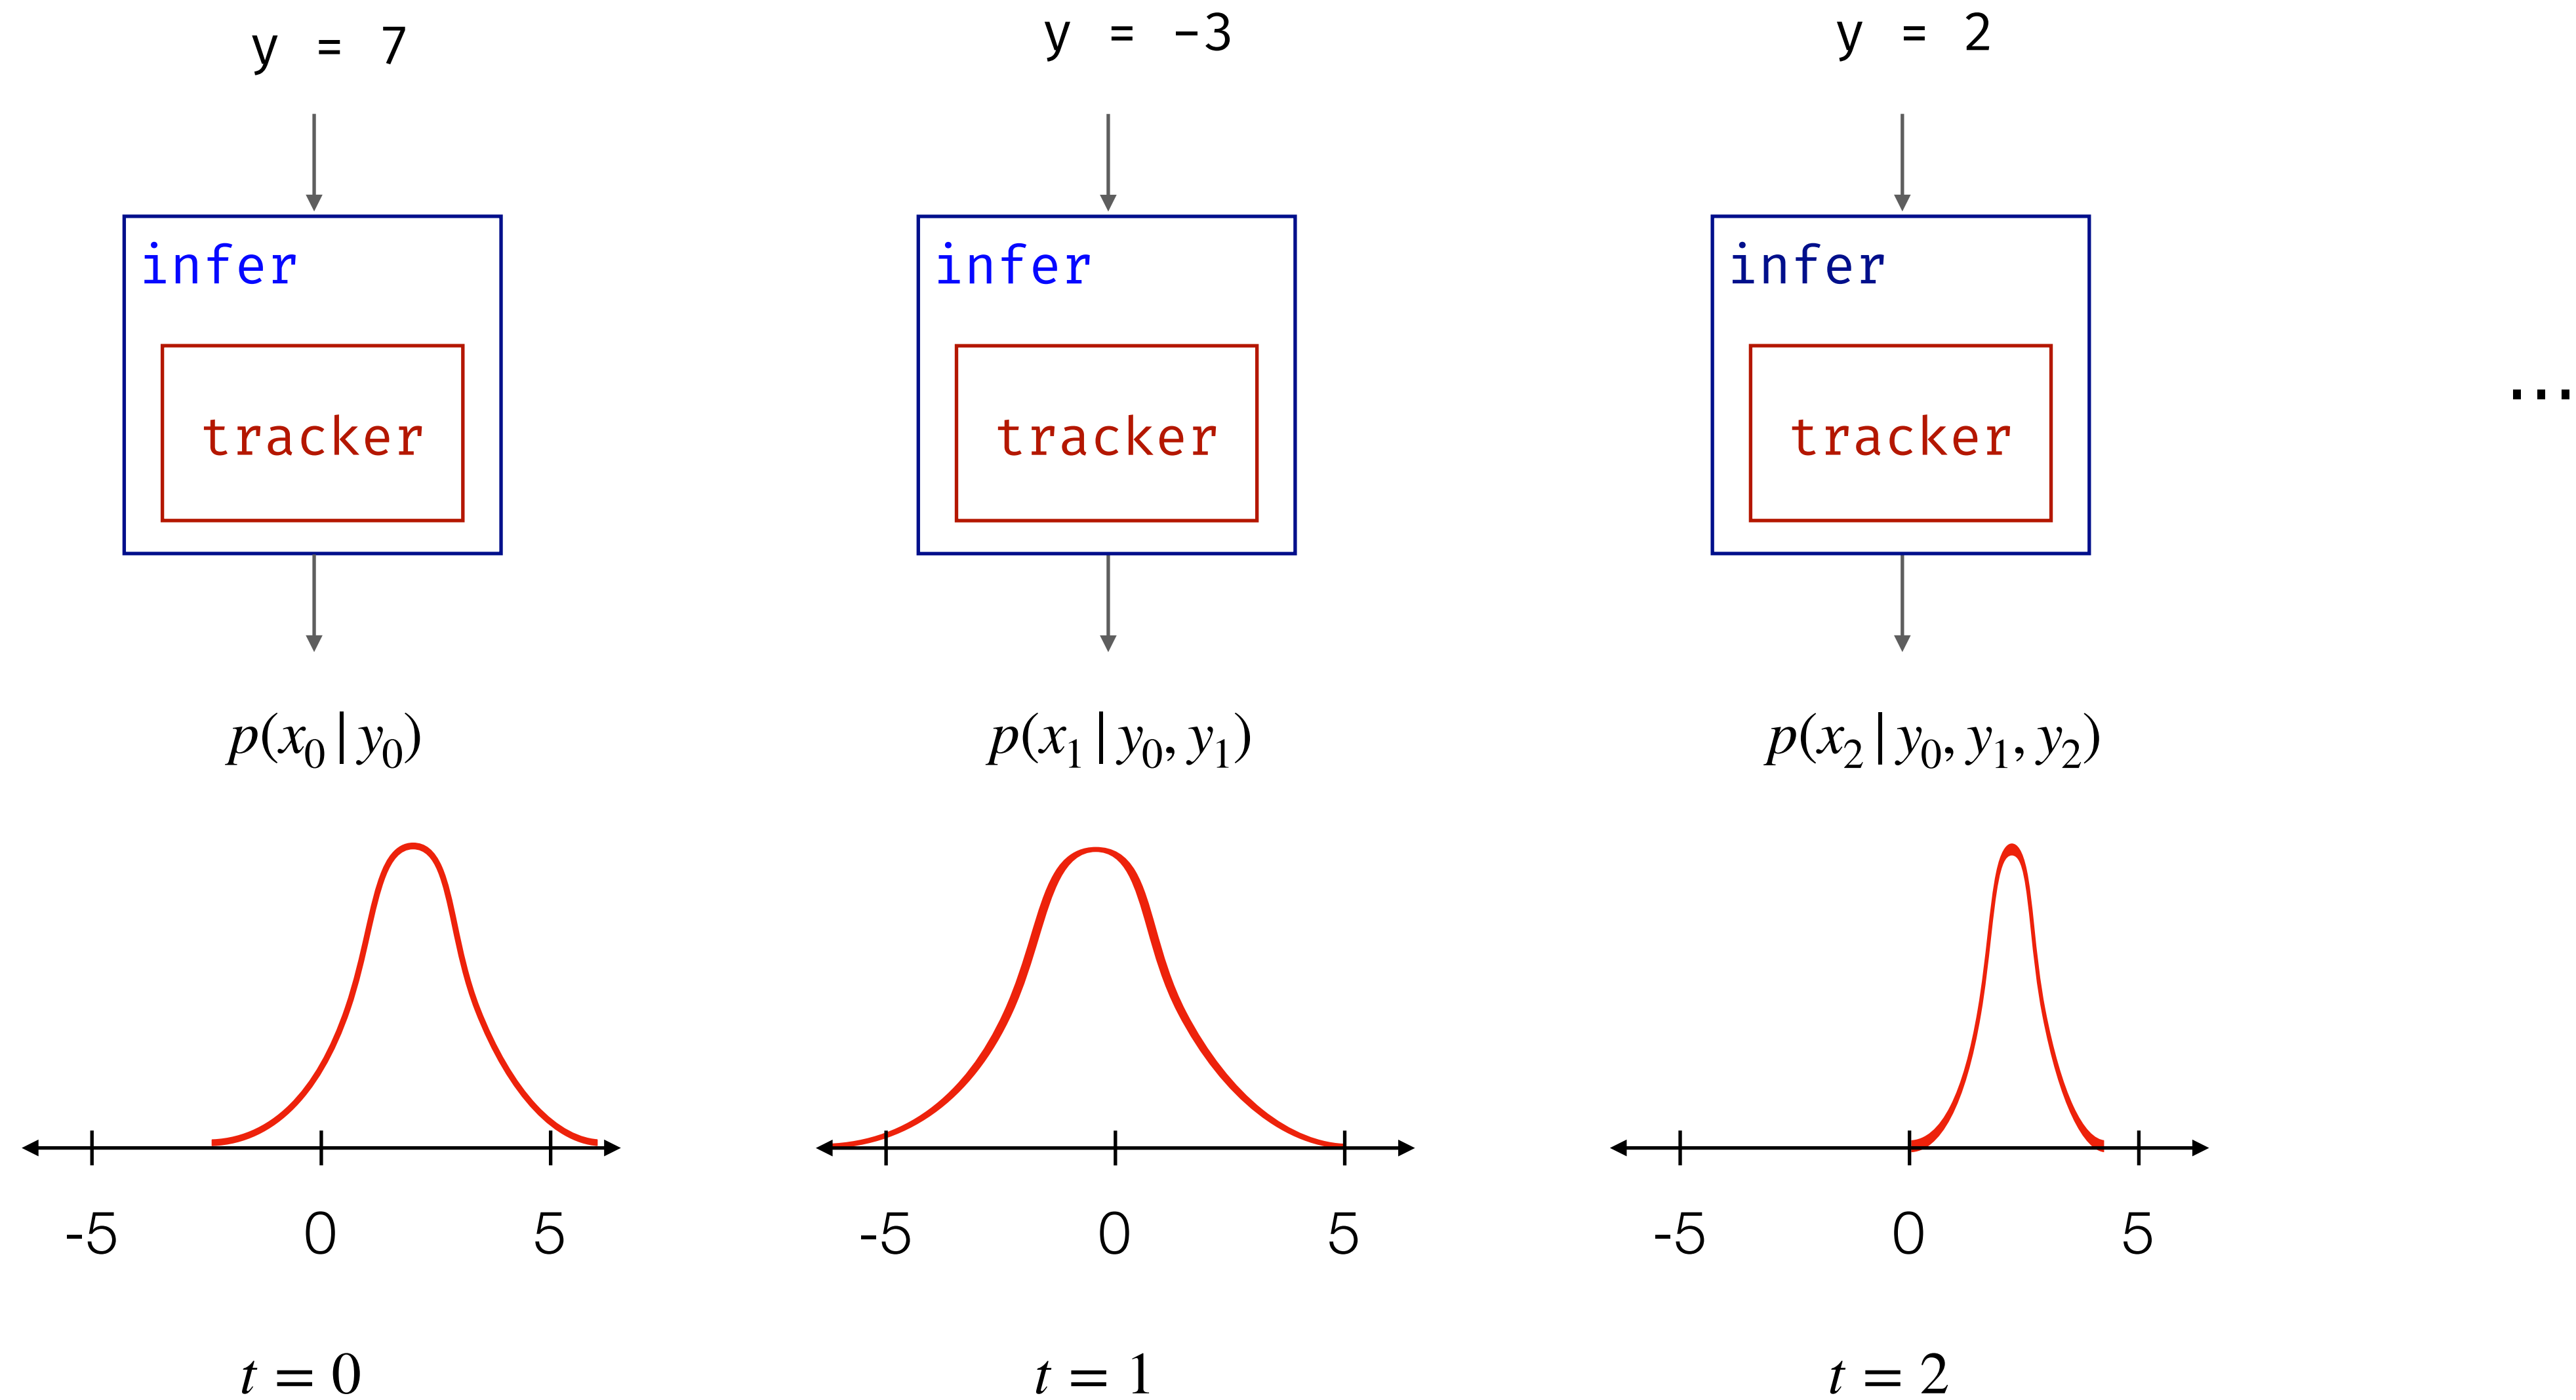
```
let proba tracker (y) = x where
  rec x = x0 → sample(mv_gaussian(f * (pre x), q))
  and () = observe(mv_gaussian(h * x, r), y)
```

## Model

- Linear motion:  $X_k \sim \mathcal{N}(F X_{k-1}, Q)$
- Observation:  $Y_k \sim \mathcal{N}(H X_k, R)$



# Reactive probabilistic programming



# Demo

# Reactive semantics

---

## Reactive Probabilistic Programming

# Deterministic streams

Initial state, transition function




$$\text{CoStream}(T, S) = S \times (S \rightarrow S \times T)$$

Constant 1:  $\text{unit} \rightarrow \text{unit} \times \text{int}$

- Initial state: the value of type unit
- Step function: return the constant 1

Unit delay **pre**  $e: (S \times T) \times (S \times T \rightarrow (S \times T) \times T)$

- Initial state: the initial state of e and default value
- Step function: the result of e is stored in the state and returned at the next iteration

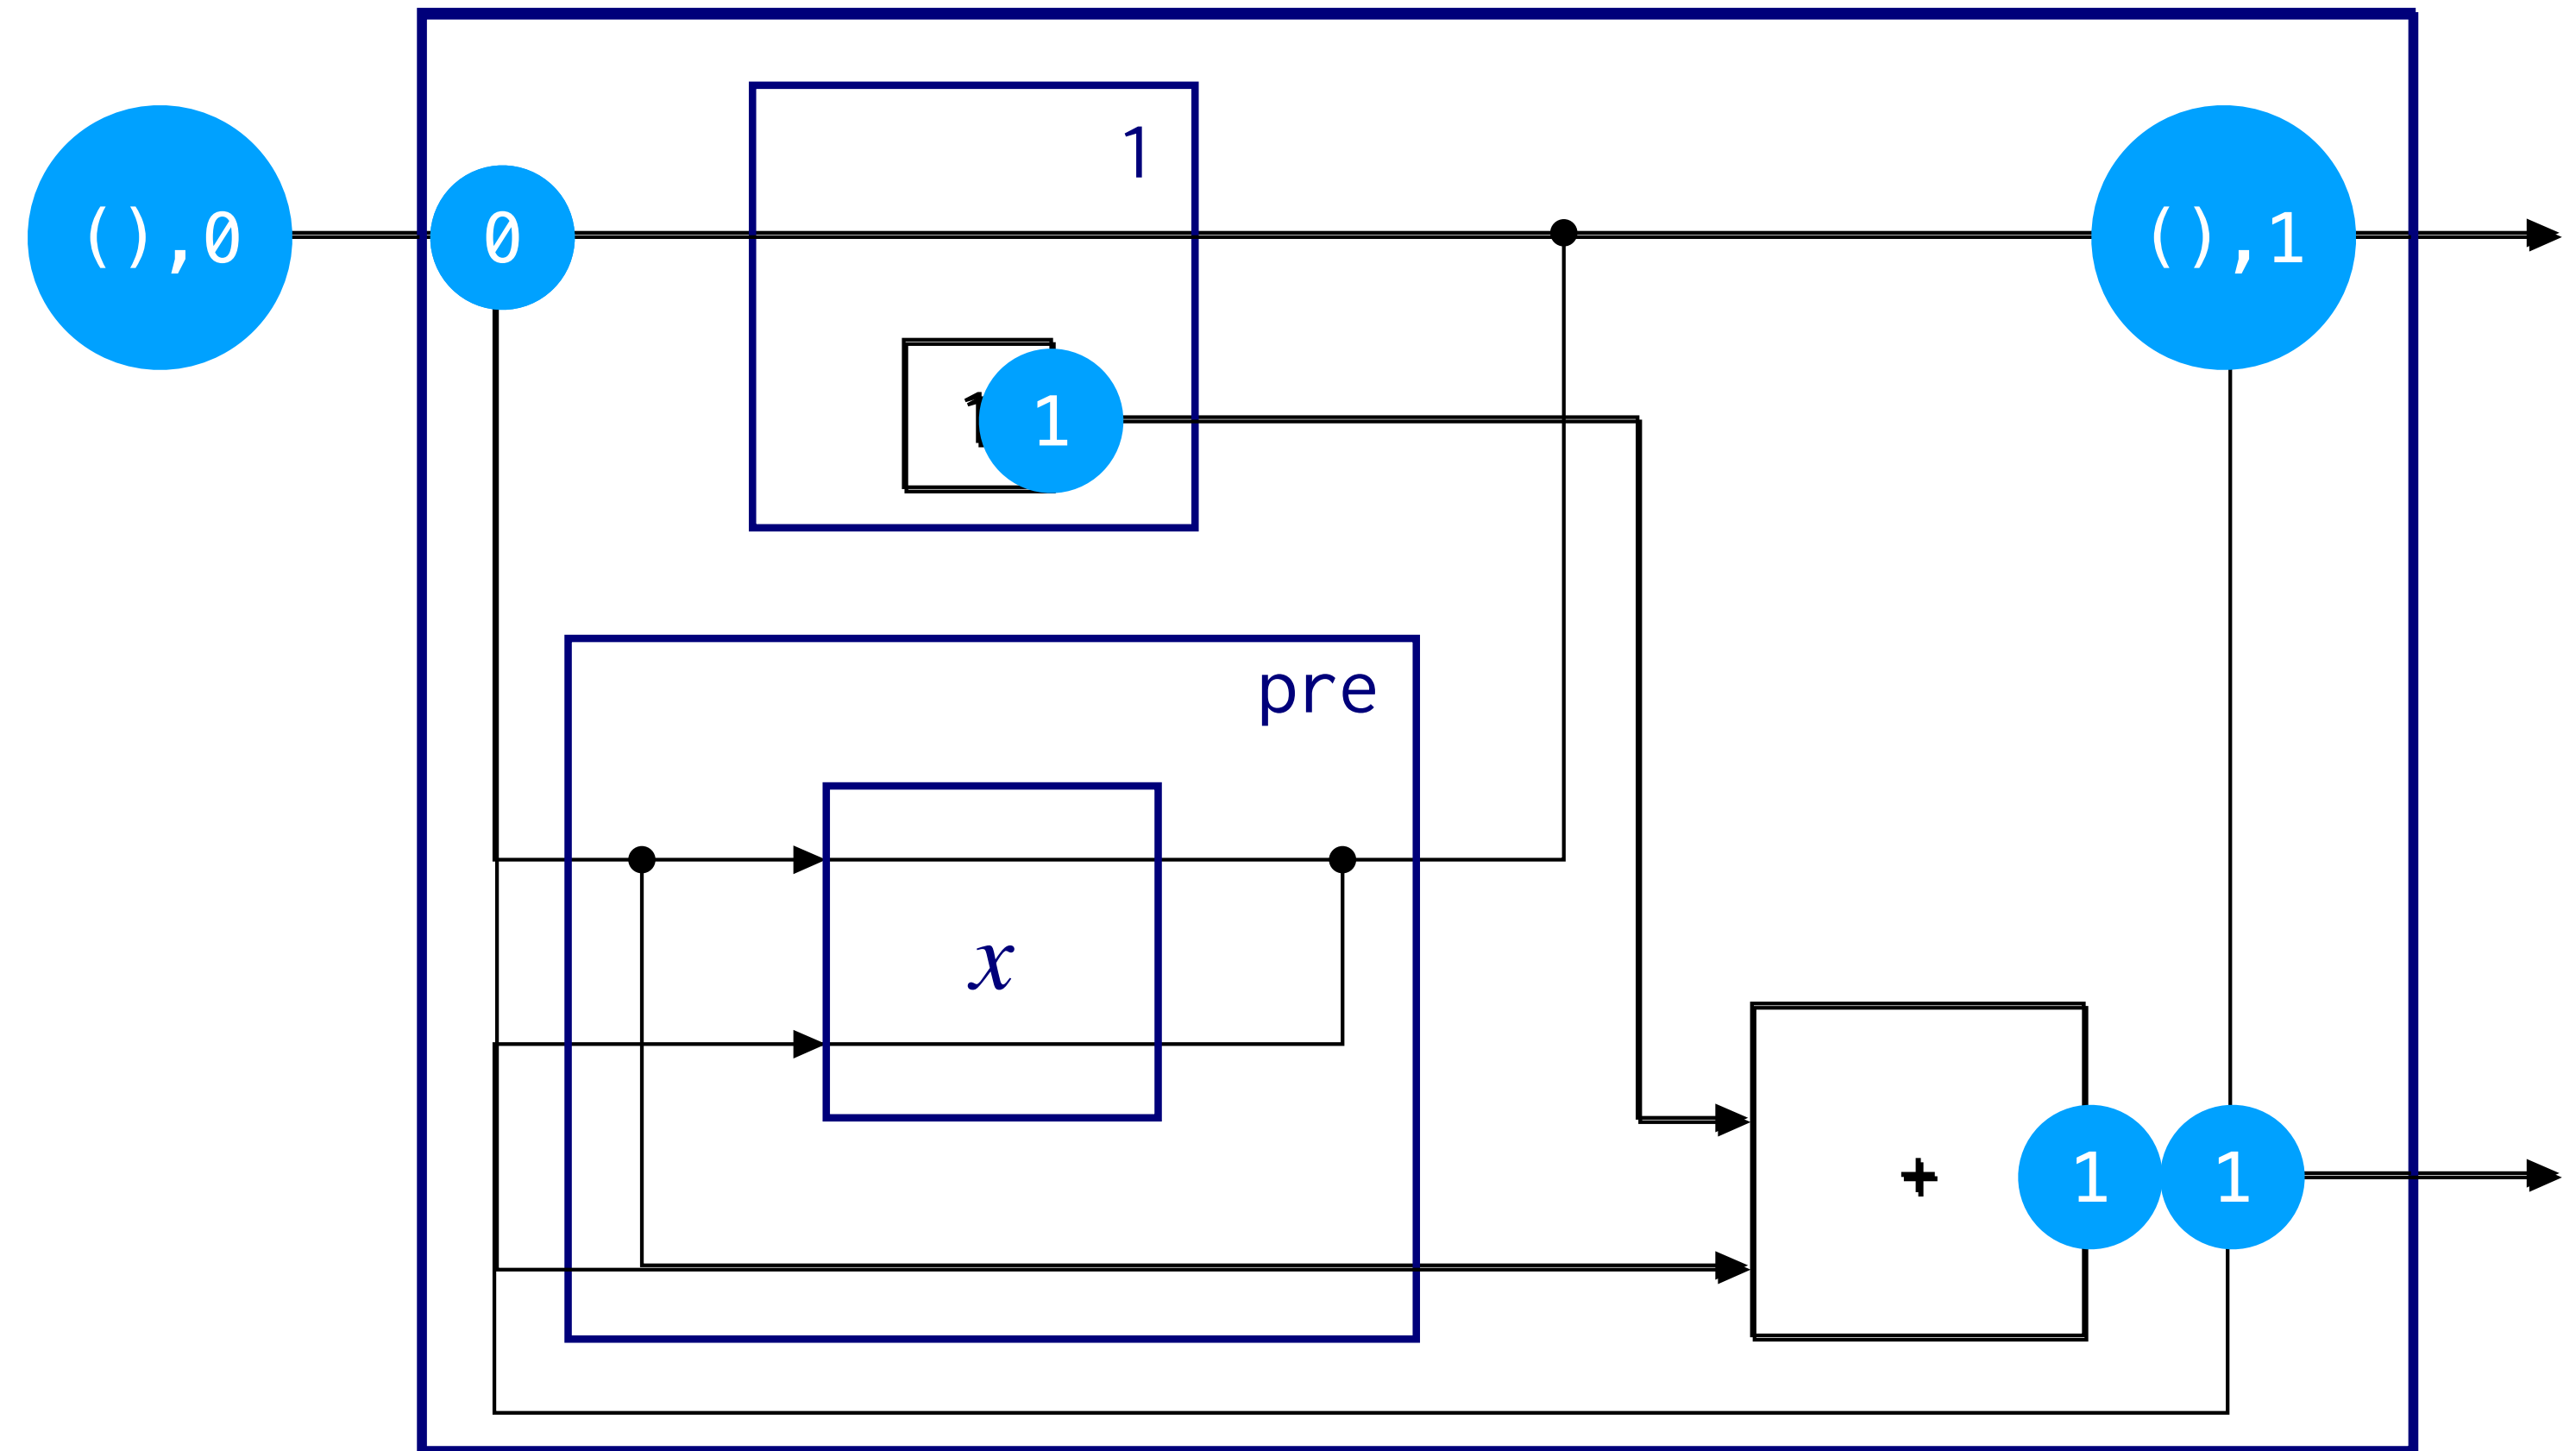


# Example

`rec`  $x = 1 + \text{pre } x$

■ Initial state:  $( ), 0$

■ Output:  $1, 2, 3, 4, \dots$

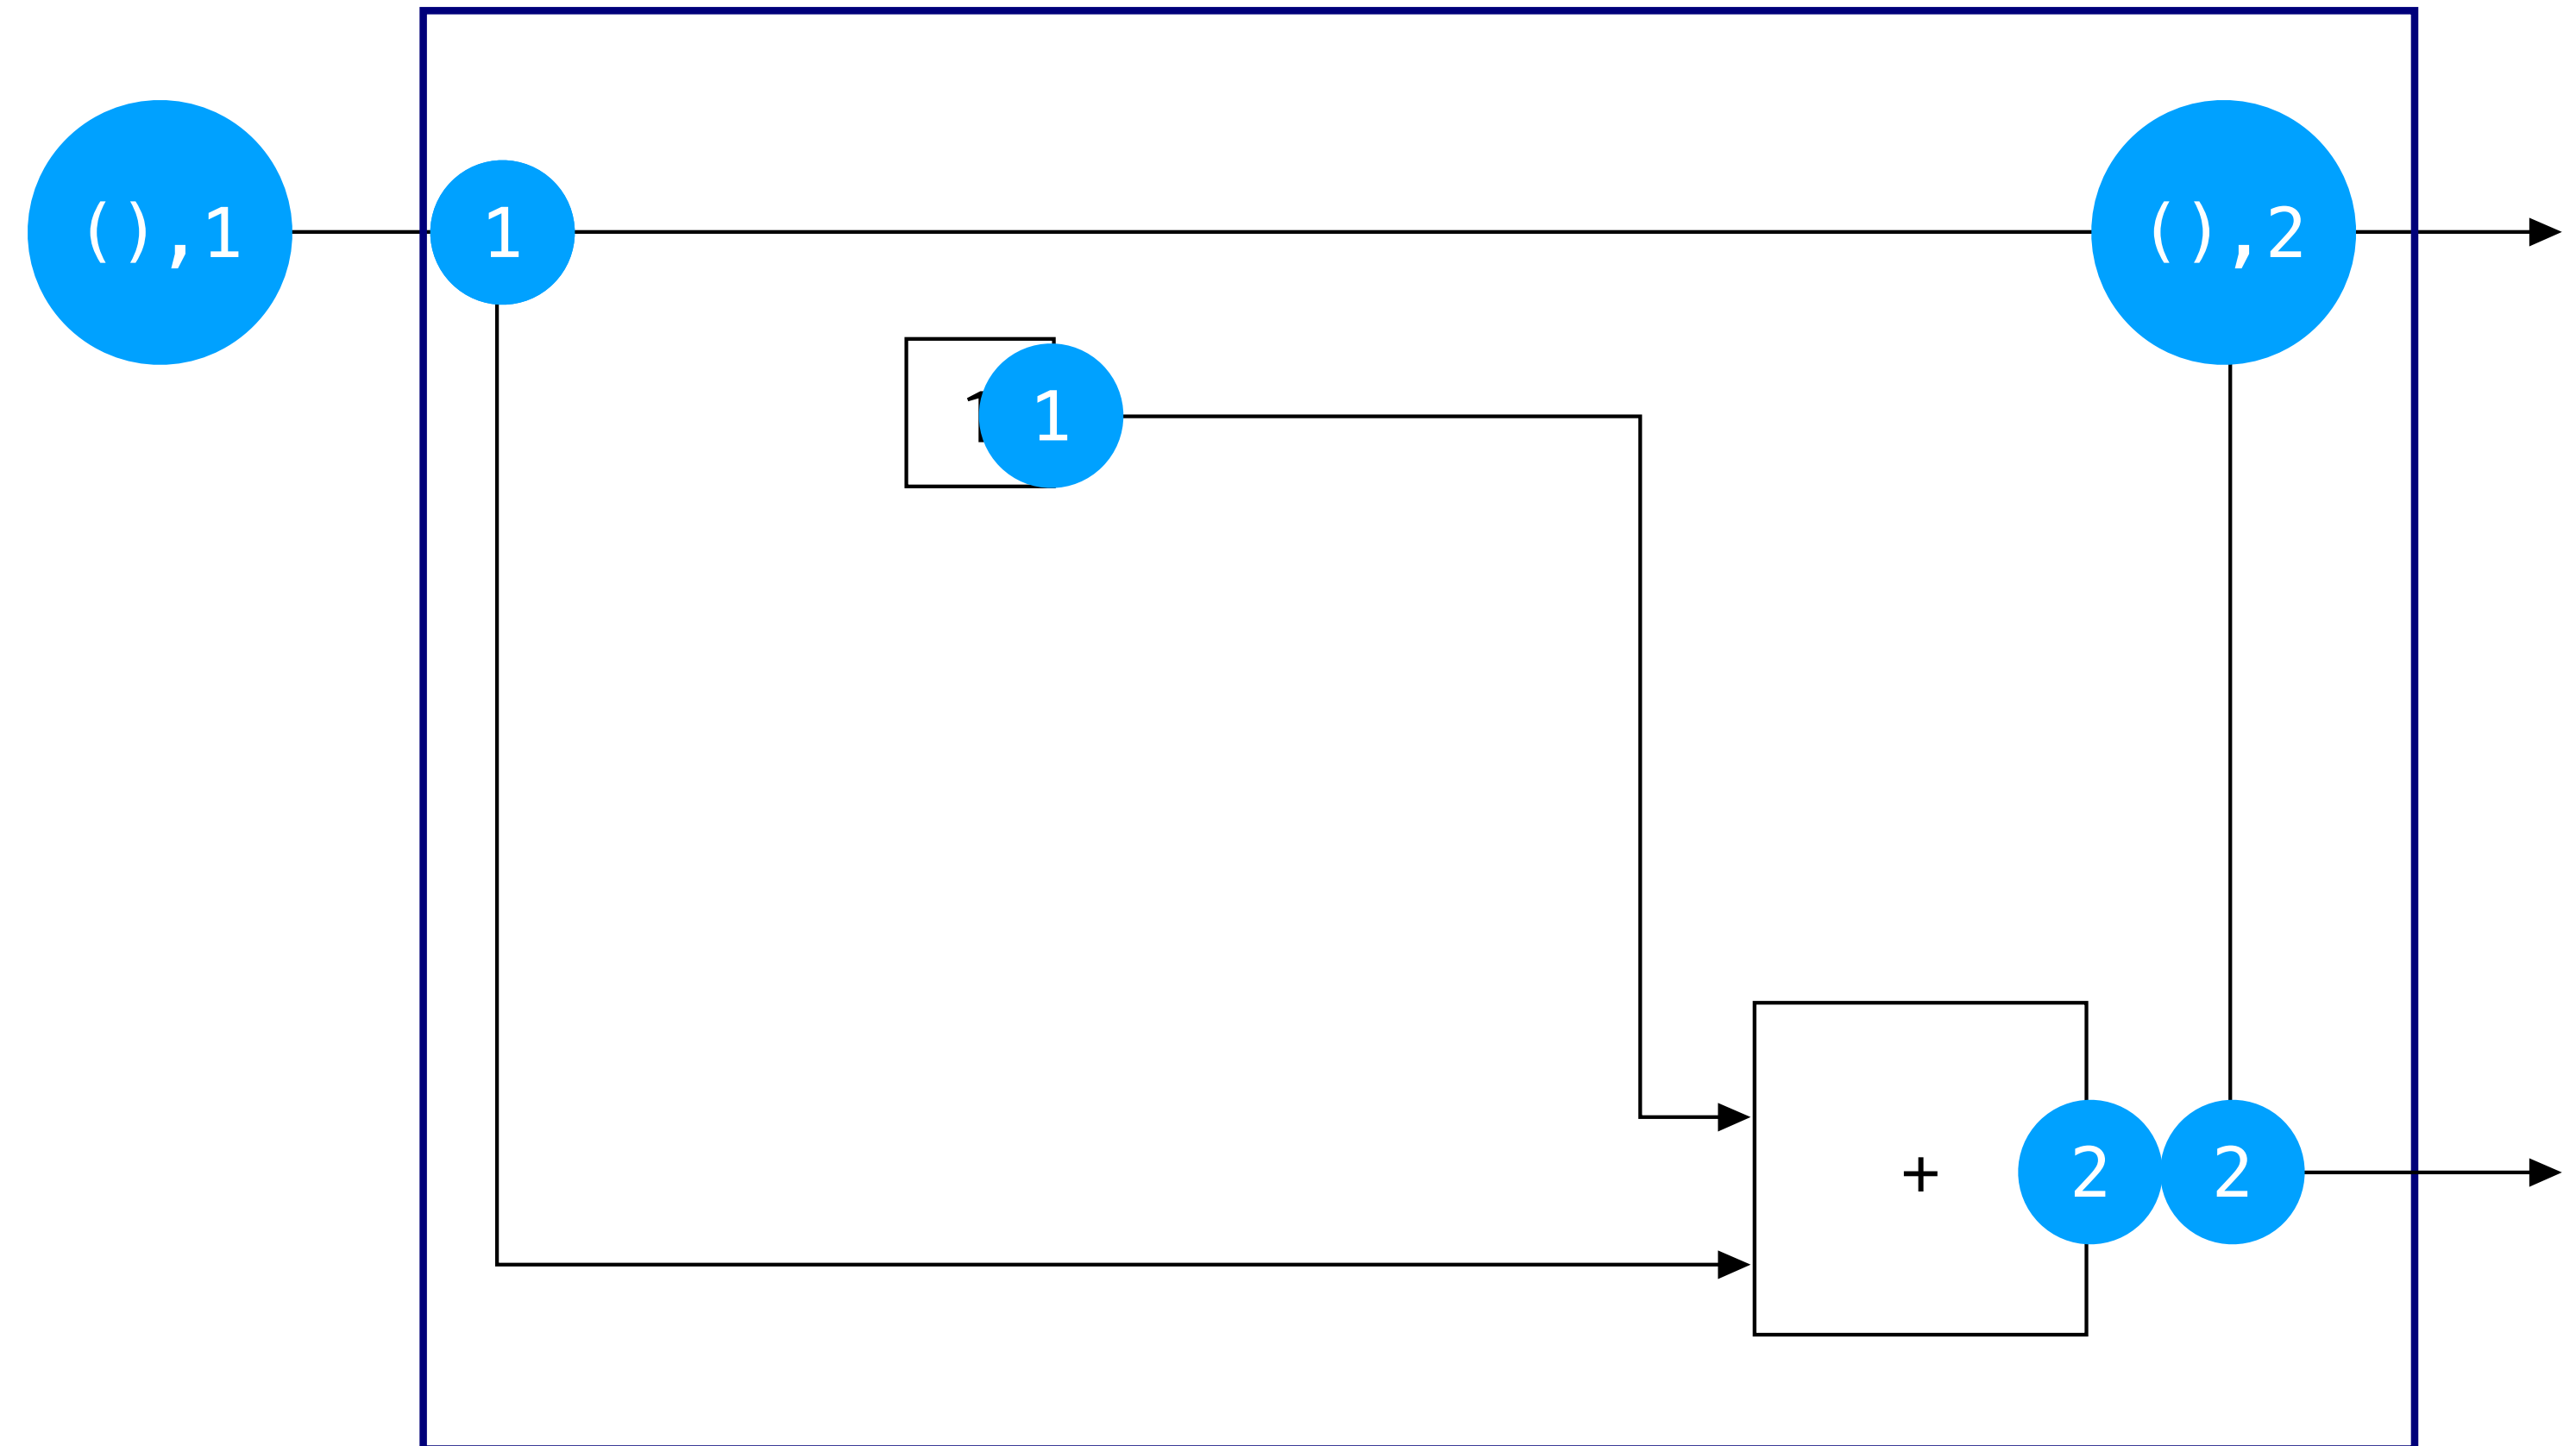


# Example



`rec x = 1 + pre x`

■ Initial state:  $()$ ,  $0$

■ Output:  $1, 2, 3, 4, \dots$



# Deterministic streams



$$\begin{aligned} \llbracket c \rrbracket_{\gamma}^{\text{init}} &= () \\ \llbracket c \rrbracket_{\gamma}^{\text{step}} (( )) &= (), c \end{aligned}$$

$$\begin{aligned} \llbracket x \rrbracket_{\gamma}^{\text{init}} &= () \\ \llbracket x \rrbracket_{\gamma}^{\text{step}} (( )) &= (), \gamma(x) \end{aligned}$$


$$\begin{aligned} \llbracket i \rightarrow \text{pre } e \rrbracket_{\gamma}^{\text{init}} &= (i, \llbracket e \rrbracket_{\gamma}^{\text{init}}) \\ \llbracket i \rightarrow \text{pre } e \rrbracket_{\gamma}^{\text{step}} (p, s) &= \text{let } s', v = \llbracket e \rrbracket_{\gamma}^{\text{step}} (s) \text{ in } (v, s'), p \end{aligned}$$

# Deterministic vs. probabilistic

## Deterministic streams

Transition function returns a pair of state and value


$$\text{CoStream}(T, S) = S \times (S \rightarrow S \times T)$$



## Probabilistic streams

Transition function returns a **measure** over (state, value)

$$\text{CoPStream}(T, S) = S \times (S \rightarrow \Sigma_{S \times T} \rightarrow [0, \infty])$$






# Probabilistic streams

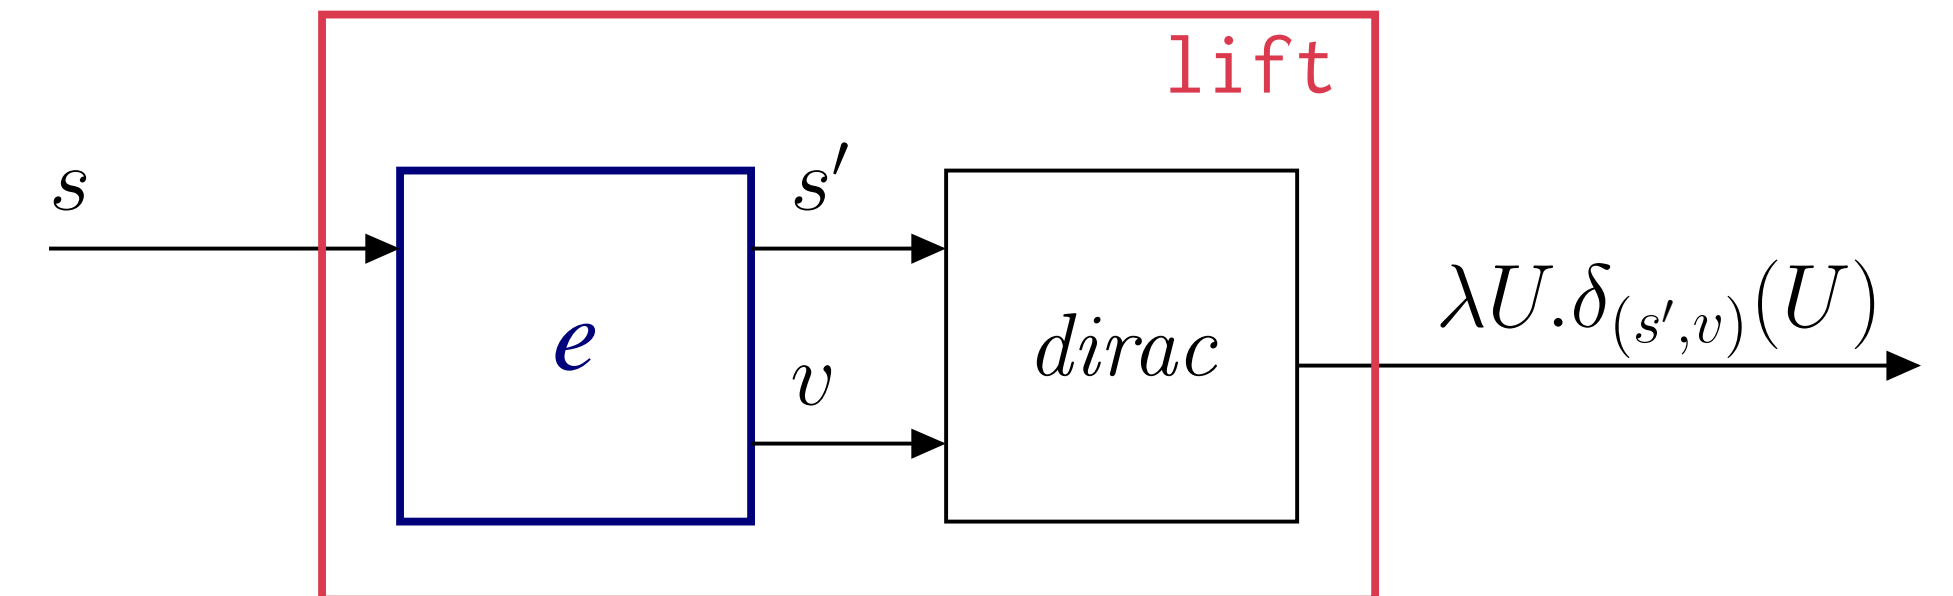
Transition function returns a *measure*

$$\text{CoPStream}(T, S) = S \times (S \rightarrow \Sigma_{S \times T} \rightarrow [0, \infty])$$



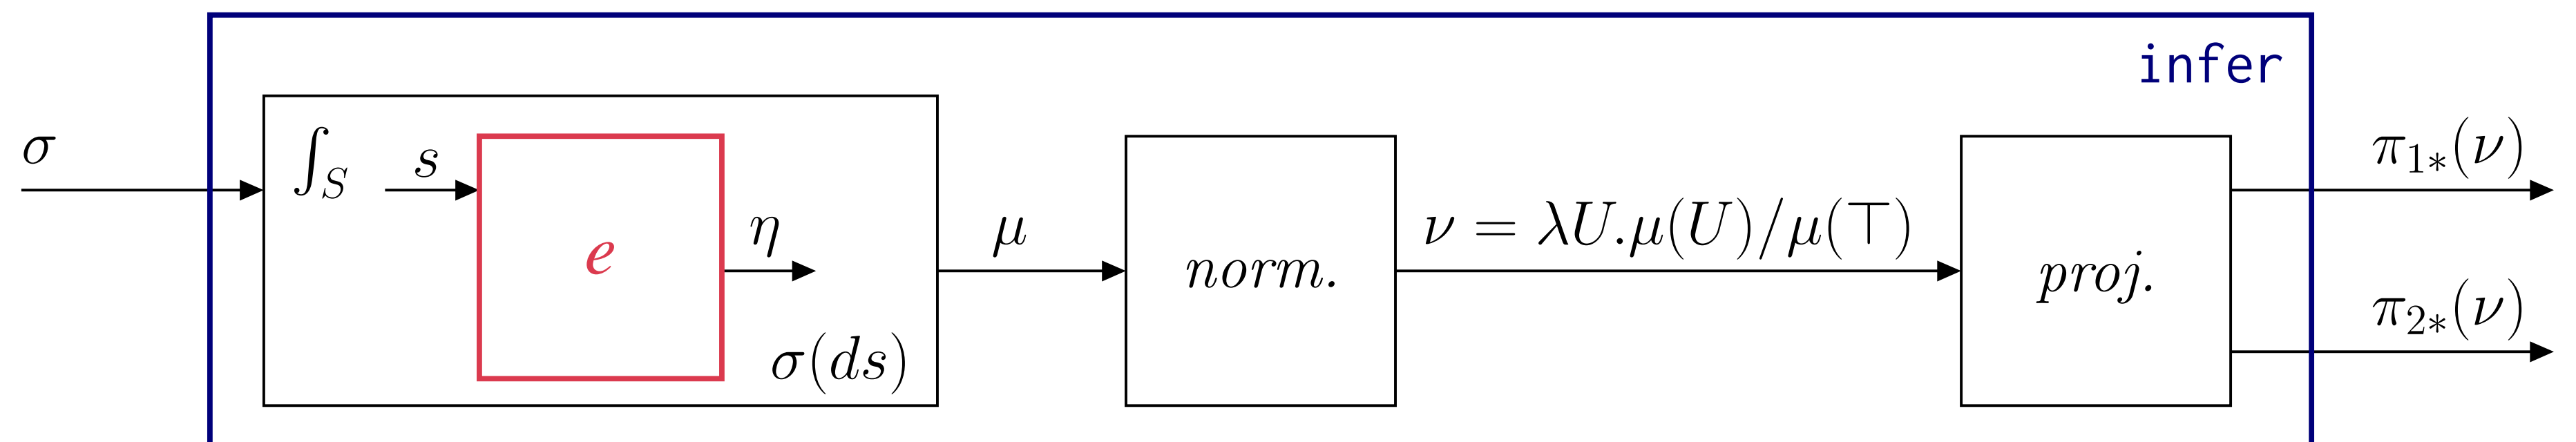
*lift* turns a deterministic expression into a probabilistic one

$$\text{lift}: S \times (S \rightarrow S \times T) \rightarrow S \times (S \rightarrow \Sigma_{S \times T} \rightarrow [0, \infty])$$


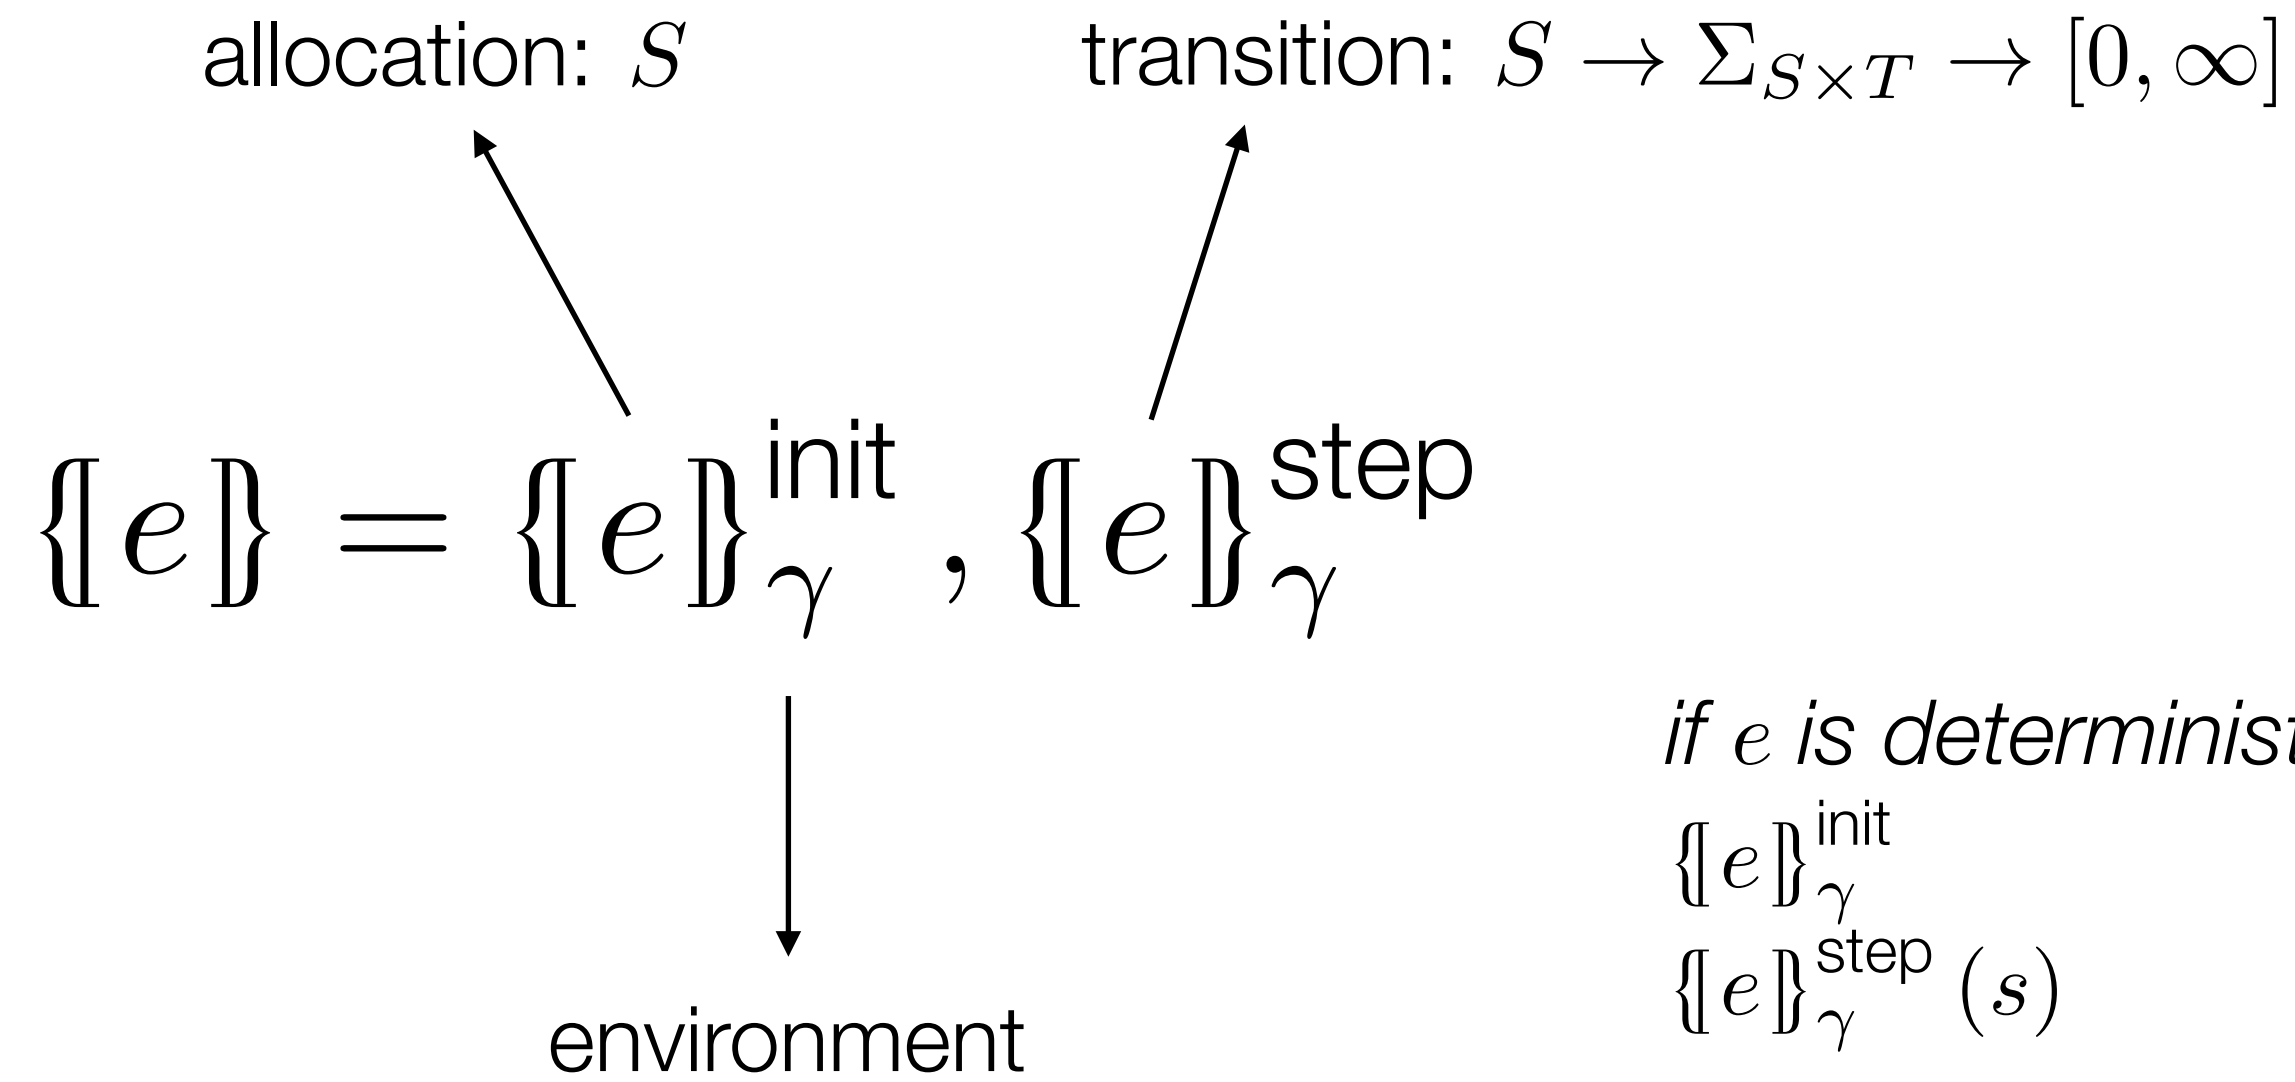


*infer* turns probabilistic expressions to a pair of distributions

$$\text{infer}: S \times (S \rightarrow \Sigma_{S \times T} \rightarrow [0, \infty]) \rightarrow S \text{ dist} \times (S \text{ dist} \rightarrow S \text{ dist} \times T \text{ dist})$$



# Probabilistic streams



*if  $e$  is deterministic*

$$\{e\}_{\gamma}^{\text{init}} \\ \{e\}_{\gamma}^{\text{step}}(s)$$

$$\{\text{sample}(e)\}_{\gamma}^{\text{init}} \\ \{\text{sample}(e)\}_{\gamma}^{\text{step}}(s)$$

$$\{\text{observe}(e_1, e_2)\}_{\gamma}^{\text{init}} \\ \{\text{observe}(e_1, e_2)\}_{\gamma}^{\text{step}}(s_1, s_2)$$

$$= \llbracket e \rrbracket_{\gamma}^{\text{init}} \\ = \text{let } s', v = \llbracket e \rrbracket_{\gamma}^{\text{step}}(s) \text{ in } \delta_{s', v}$$

$$= \llbracket e \rrbracket_{\gamma}^{\text{init}} \\ = \text{let } s', \mu = \llbracket e \rrbracket_{\gamma}^{\text{step}}(s) \text{ in } \int \mu(dv) \delta_{s', v}$$

$$= \llbracket e_1 \rrbracket_{\gamma}^{\text{init}}, \llbracket e_2 \rrbracket_{\gamma}^{\text{init}} \\ = \text{let } s'_1, \mu = \llbracket e_1 \rrbracket_{\gamma}^{\text{step}}(s_1) \text{ in} \\ \text{let } s'_2, v = \llbracket e_2 \rrbracket_{\gamma}^{\text{step}}(s_2) \text{ in} \\ \mu_{\text{pdf}}(v) * \delta_{(s'_1, s'_2), ()}$$

# Streaming inference

---

Reactive Probabilistic Programming

# Particle filter

```
let proba tracker (y) = x where  
  rec x = sample (gaussian (0, 10)) → gaussian (pre x, 1)  
  and () = observe (gaussian (x, 1), y)
```

$t = 0$


```
sample (gaussian (0, 10))  
observe (gaussian (x, 1), 3)
```

$t = 1$

```
sample (gaussian (pre x, 1))  
observe (gaussian (x, 1), 5)
```

$t = 2$

```
sample (gaussian (pre x, 1))  
observe (gaussian (x, 1), ...)
```




# Delayed sampling

```
let proba tracker (y) = x where
  rec x = sample (gaussian (0, 10)) → gaussian (pre x, 1))
  and () = observe (gaussian (x, 1), y)
```

$t = 0$

```
sample (gaussian (0, 10))
observe (gaussian (x, 1), 3)
```




$t = 1$

```
sample (gaussian (pre x, 1))
observe (gaussian (x, 1), 5)
```

$t = 2$

Unbounded resources



# Benchmarks

MODEL	PF		DS		SSI	
	# PART.	TIME (MS)	# PART.	TIME (MS)	# PART.	TIME (MS)
Beta-Bernoulli	200	23.05 (22.54-23.90)	✓ 1	0.28 (0.27-0.28)	✓ 1	0.65 (0.65-0.66)
Gaussian-Gaussian	3000	877.44 (689.04-890.68)	150	62.99 (61.69-65.86)	150	182.57 (181.40-183.54)
Kalman-1D	15	3.27 (3.26-3.28)	✓ 1	0.33 (0.33-0.34)	✓ 1	1.15 (1.14-1.15)
Outlier	700	222.27 (220.76-223.84)	65	43.81 (43.45-46.48)	65	125.88 (125.27-128.08)
Robot	85	771.32 (767.98-775.51)	✓ 1	91.44 (90.94-92.07)	✓ 1	96.40 (96.21-97.53)
SLAM		✗	800	2812.55 (2755.99-2853.89)	800	5649.30 (5619.59-5675.81)
MTT		✗	60	2889.11 (2615.76-3244.30)	60	4457.79 (4068.35-4996.20)
Tree	150	35.55 (35.41-35.68)	90	58.83 (58.55-59.74)	✓ 1	2.67 (2.66-2.70)
Wheels	550	246.48 (245.06-248.75)	550	699.12 (672.25-713.64)	✓ 1	8.04 (8.00-8.10)
Delayed GPS	150	1221.00 (1218.76-1230.67)	9	304.73 (303.17-306.31)	✓ 1	108.55 (108.02-109.07)

Time to reach a target accuracy

Baseline:  
SDS with 1,000 particles

SSI computations cost more,  
but can outperform DS

Trade-off: cost of symbolic computations vs. precision


# Static Analysis

---

Reactive Probabilistic Programming

# Bounded Memory Delayed Sampling?

```
proba tracker (y) = x where  
  rec x = sample (gaussian (0, 10) → gaussian (pre x, 1))  
  and () = observe (gaussian (x, 1), y)
```




Yes!



# Bounded Memory Delayed Sampling?

```
proba tracker (y) = x, x0 where
  rec init x0 = sample (gaussian (0, 10))
  and x = x0 → sample (gaussian (pre x, 1))
  and () = observe (gaussian (x, 1), y)
```

Can we determine if a given program will run in bounded memory?



No!

# Trace: Abstract Execution

```
proba tracker (y) = x where
  rec x = sample (gaussian (0, 10) → gaussian (pre x, 1))
  and () = observe (gaussian (x, 1), y)
```

	trace	state	time
random variable →	$x_0 \leftarrow \perp ::$ $y_0 \leftarrow x_0 ::$	$X = x_0$	$t = 0$
observation →	$\text{observe } y_0 ::$		
	$x_1 \leftarrow x_0 ::$ $y_1 \leftarrow x_1 ::$ $\text{observe } y_1 ::$	$x = x_1, \text{ pre } x = x_0$	$t = 1$
	$x_2 \leftarrow x_1 ::$ $y_2 \leftarrow x_2 ::$ ...	$x = x_2, \text{ pre } x = x_1$	$t = 2$

# Static Analysis for Delayed Sampling

## Semantic properties

### m-consumed property

Chains of variables before an observe are bounded

### unseparated paths property

Chains of variables referenced in the state are bounded

Theorem: *The program satisfies these two properties iff it executes in bounded memory*

---

## Static analysis

Track variables introduced but not used yet

Track maximal path between pairs of variable in the state

Theorem: *The program pass the analysis if it executes in bounded memory*

# *m*-consumed Property

```
proba tracker (y) = x where
  rec x = sample (gaussian (0, 10) → gauss
  and () = observe (gaussian (x, 1), y)
```

Chains of variables before an observe are bounded

	trace	state	time
	$x_0 \leftarrow \perp ::$	$x = x_0$	$t = 0$
$x_0$ is 1-consumed →	$y_0 \leftarrow x_0 ::$		
$y_0$ is 0-consumed →	observe $y_0 ::$		
	$x_1 \leftarrow x_0 ::$	$x = x_1, \text{ pre } x = x_0$	$t = 1$
$x_1$ is 1-consumed →	$y_1 \leftarrow x_1 ::$		
$y_1$ is 0-consumed →	observe $y_1 ::$		
	$x_2 \leftarrow x_1 ::$	$x = x_2, \text{ pre } x = x_1$	$t = 2$
	$y_2 \leftarrow x_2 ::$		
	...		

Yes!

# Unseparated Paths Property

```
proba tracker (y) = x where
  rec x = sample (gaussian (0, 10) → gauss
  and () = observe (gaussian (x, 1), y)
```

Chains of variables referenced in the state are bounded

trace	state	time
$x_0 \leftarrow \perp ::$ $y_0 \leftarrow x_0 ::$ observe $y_0 ::$	$x = x_0$	$t = 0$
$x_1 \leftarrow x_0 ::$ $y_1 \leftarrow x_1 ::$ observe $y_1 ::$	$x = x_1, \text{ pre } x = x_0$	$t = 1$
$x_2 \leftarrow x_1 ::$ $y_2 \leftarrow x_2 ::$ ...	$x = x_2, \text{ pre } x = x_1$	$t = 2$

Yes!

# Unseparated Paths Property

```
proba tracker (obs) = x where
  rec init x0 = sample (gaussian (0, 10))
  and x = x0 → sample (gaussian (pre x, 1))
  and () = observe (gaussian (x, 1), y)
```

Chains of variables referenced in the state are bounded

trace	state	time
$x_0 \leftarrow \perp ::$	$x = x_0$	$t = 0$
$y_0 \leftarrow x_0 ::$	$x_0 = x_0$	
$\text{observe } y_0 ::$		
$x_1 \leftarrow x_0 ::$	$x = x_1, \text{ pre } x = x_0$	$t = 1$
$y_1 \leftarrow x_1 ::$	$x_0 = x_0$	
$\text{observe } y_1 ::$		
$x_2 \leftarrow x_1 ::$	$x = x_2, \text{ pre } x = x_1$	$t = 2$
$y_2 \leftarrow x_2 ::$	$x_0 = x_0$	
...		

No!

# References

**Delayed Sampling and Automatic Rao-Blackwellization of Probabilistic Programs.**

Lawrence Murray, Daniel Lundén, Jan Kudlicka, David Broman, Thomas B. Schön  
AISTATS 2017

**Semi-Symbolic Inference for Efficient Streaming Probabilistic Programming**

Eric Atkinson, Charles Yuan, Guillaume Baudart, Louis Mandel, Michael Carbin  
OOPSLA 2022

**Reactive probabilistic programming**

Guillaume Baudart, Louis Mandel, Eric Atkinson, Benjamin Sherman, Marc Pouzet, Michael Carbin  
PLDI 2020

<https://github.com/IBM/probzelus>

**Statically Bounded-Memory Delayed Sampling for Probabilistic Streams**

Eric Atkinson, Guillaume Baudart, Louis Mandel, Charles Yuan, Michael Carbin  
OOPLSA 2021